Advanced Computer Vision

Elisa Ricci & Nicu Sebe

some slides from S. Hochreiter, I. Laptev, G. Mori, M. Ryoo

Top 10 AI technology trends

Deep learning theory

The information bottleneck principle explains how a deep neural network learns.

Capsule networks

New type of deep neural network that learns with fewer errors and less data, by preserving key hierarchical relationships.

Deep reinforcement learning

This technique combines reinforcement learning with deep neural networks to learn by interacting with the environment.

Generative adversarial networks

A type of unsupervised deep learning system, implemented as two competing neural networks, enabling machine learning with less human intervention.

Lean and augmented data learning

Different techniques that enable a model to learn from less data or synthetic data.

Probabilistic programming

A high-level language that makes it easy for developers to define probability models.

Hybrid learning models

Approach that combines different types of deep neural networks with probabilistic approaches to model uncertainty.

Automated machine learning

Technique for automating the standard workflow of machine learning.

Digital twin

A virtual model used to facilitate detailed analysis and monitoring of physical or psychological systems.

Explainable artificial intelligence

Machine learning techniques that produce more explainable models while maintaining high performance.

Successes of Deep Learning in AI

The New York Times

A Learning Advance in Artificial Intelligence Rivals Human Abilities

FLOWER
FIELD
SKY
CLOUDS

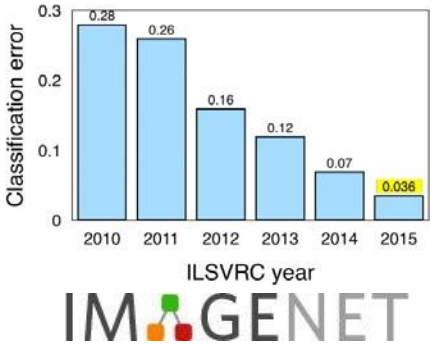

Deep Learning for self-driving cars

Google's DeepMind Masters Atari Games

Google Translate

English | Russian | Chinese (Simplified)

Time flies like an arrow

时间过得很快像箭

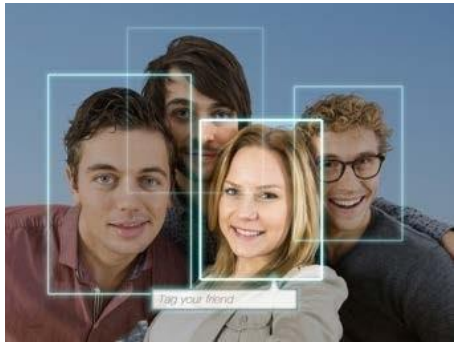

Face Recognition

So is AI solved?

pedestrian detection FAIL (Volvo S60 Pedestrian Detection System Test)

https://www.youtube.com/watch?v=w2pwxv8rFkU

Pedestrian detection is only active at <22mph and >2mph. Once the brakes have been engaged, they do release afterwards assuming the driver who wasn't initially paying attention is now fully aware of the audible alarm and flashing red lights on the windshield. This is why this guy almost got run over: the system was working as designed but he was not behind the wheel to take back!!!

Challenges

What does computer vision involve?

Detection: are there people?

Objects and scenes: where are they?

Action recognition: what are they doing?

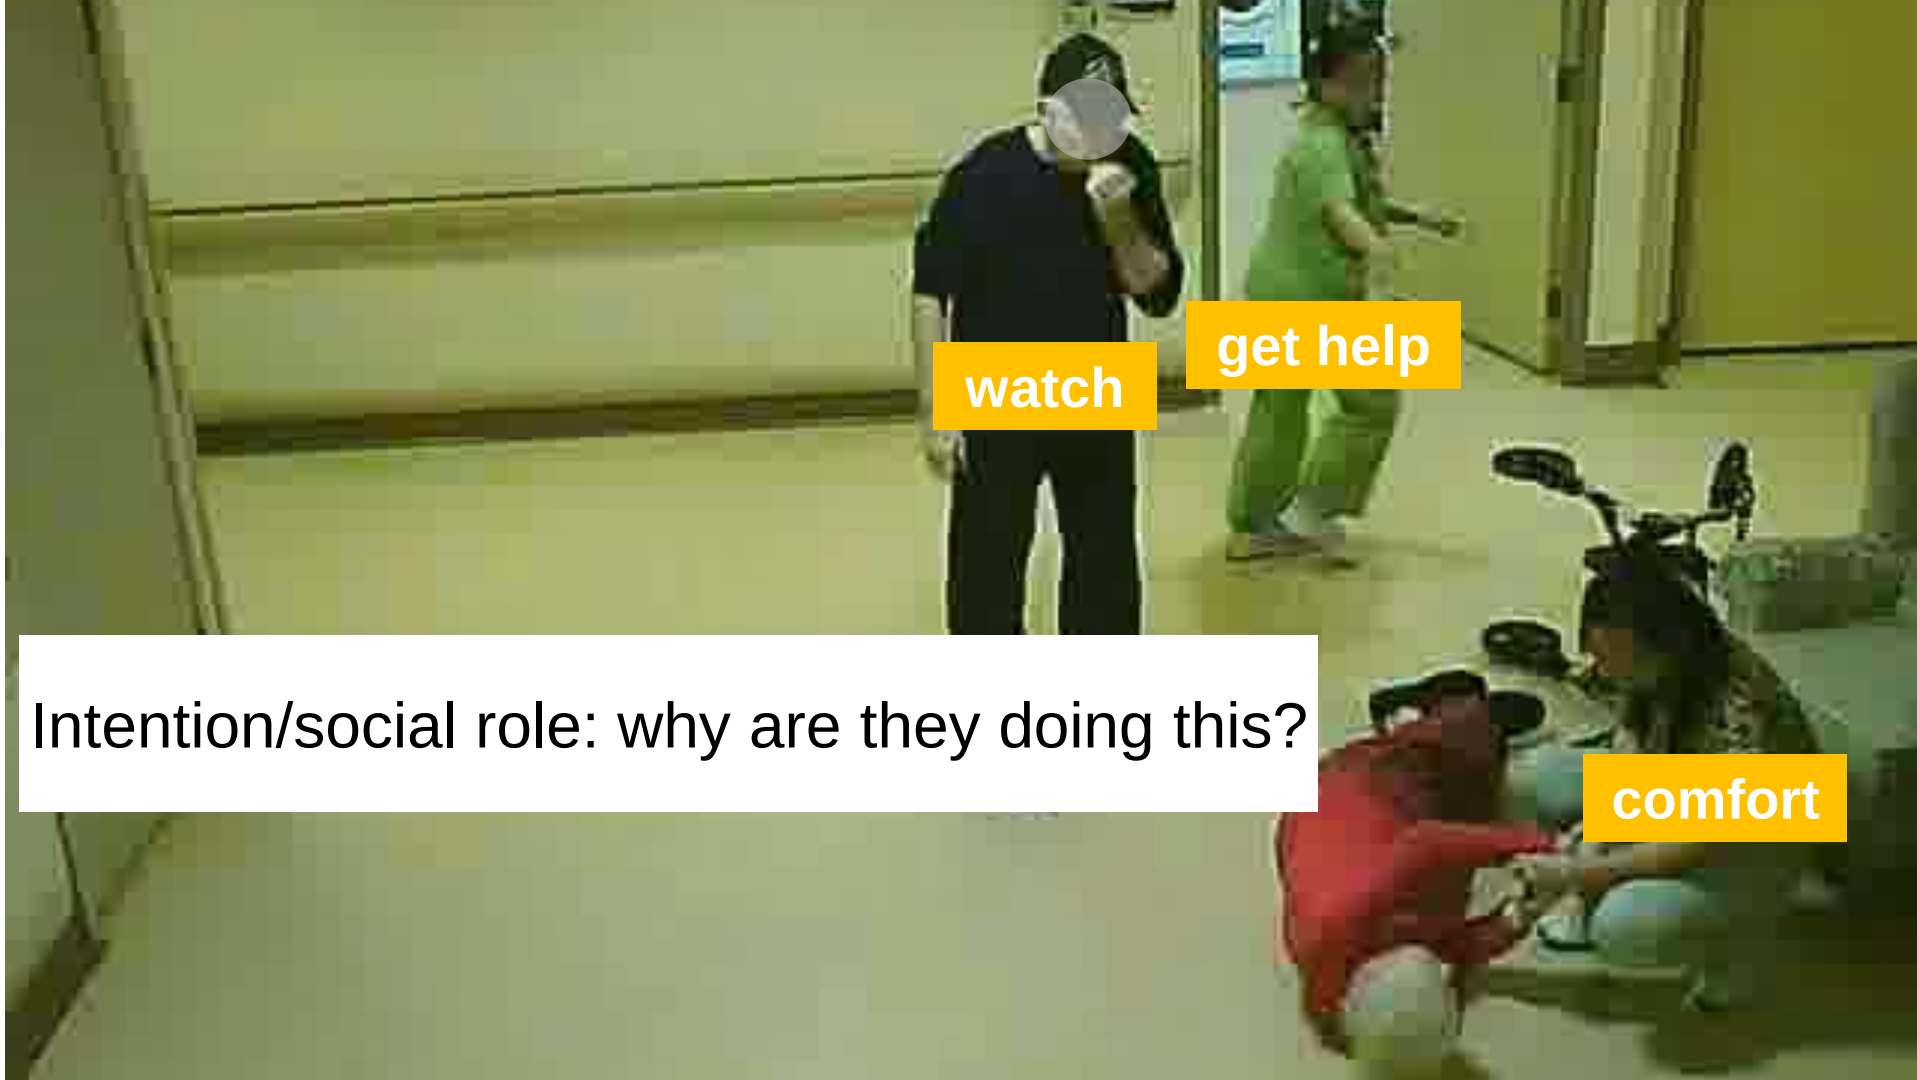

Intention/social role: why are they doing this?

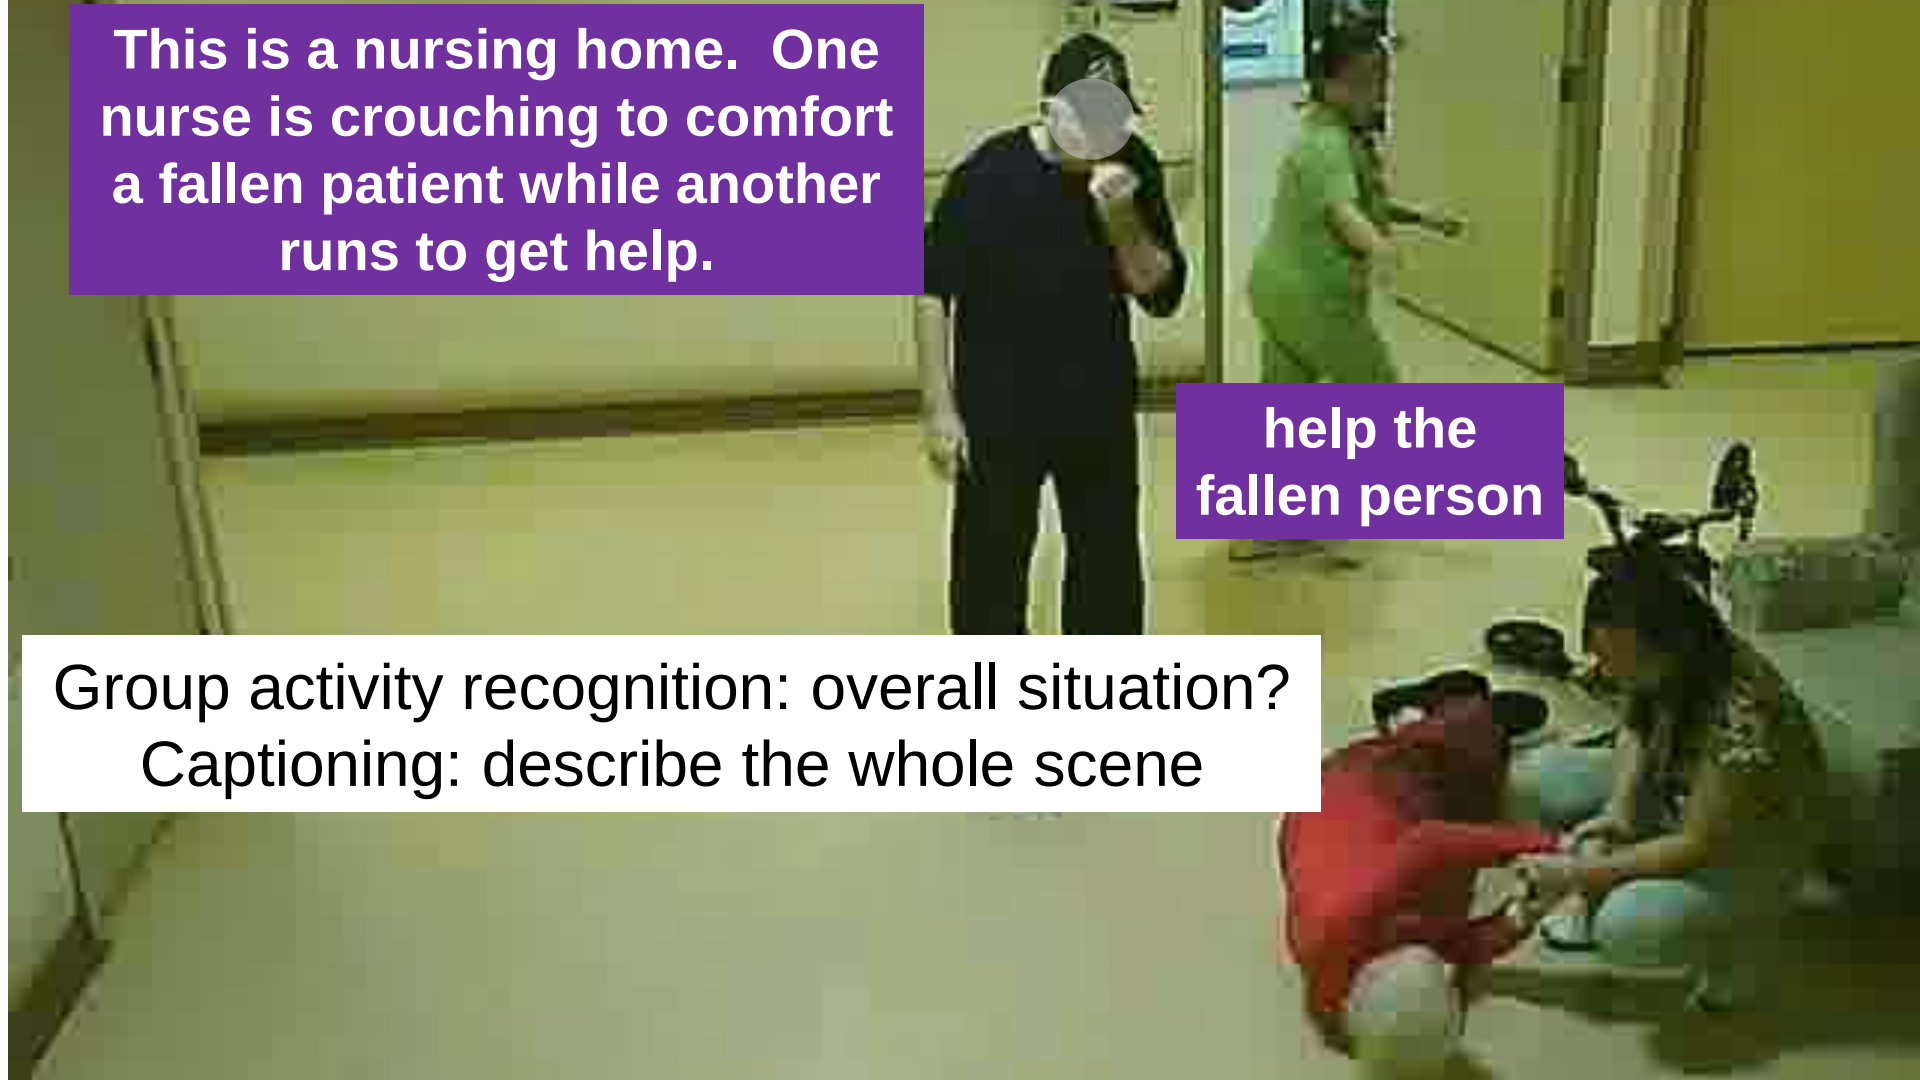

Group activity recognition: overall situation?
Captioning: describe the whole scene

"AI-complete": full semantic understanding necessary for success

Why is computer vision important?

User videos

~300 hours of videos per minute

- Video indexing and retrieval

Monitoring cameras

Streaming videos 24/7

- Surveillance
- Patient/elderly monitoring

Media

Content analysis, experience enrichment

- Recommendation systems
- Advertising
- Sports analytics

Wearables/robots

Streaming videos to be analyzed in real-time

- Lifelogging
- Robot operations and actions

Why is it difficult?

- **Large variations in appearance:** occlusions, non-rigid motion, view-point changes, clothing, etc.

Action *Hugging*:

…

- **Manual collection of training samples is prohibitive:** many classes, rare occurrence

…

- **Vocabulary is not well-defined**

Action *Open*:

…

[slide credit: Ivan Laptev]

Why is it difficult?

Humans Outperform Computers

- Humans generalize better that deep learning systems
- Humans can handle out of distribution data
- Humans are more robust to domain shifts
- Humans can handle new situations, objects, environments
- Humans are better in building abstractions

[slide credit: Sepp Hochreiter]

Example: Human activity recognition

- There are various types/levels of activities
  - The ultimate goal is to make computers recognize all of them reliably

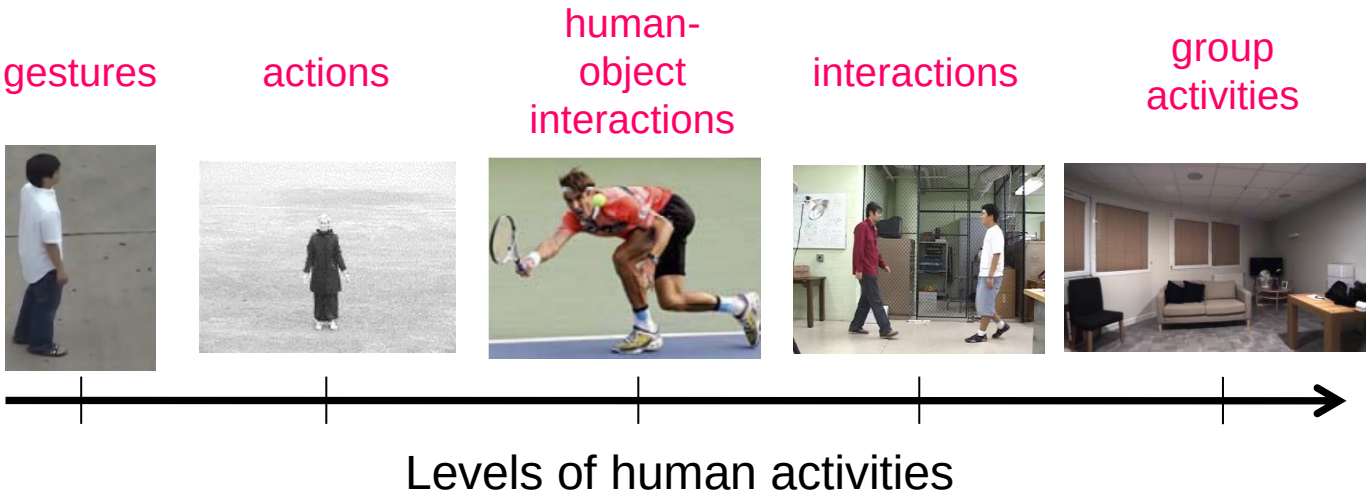

Example: Activity classification

- Categorization of segmented videos
  - Input = a video segment containing only one activity

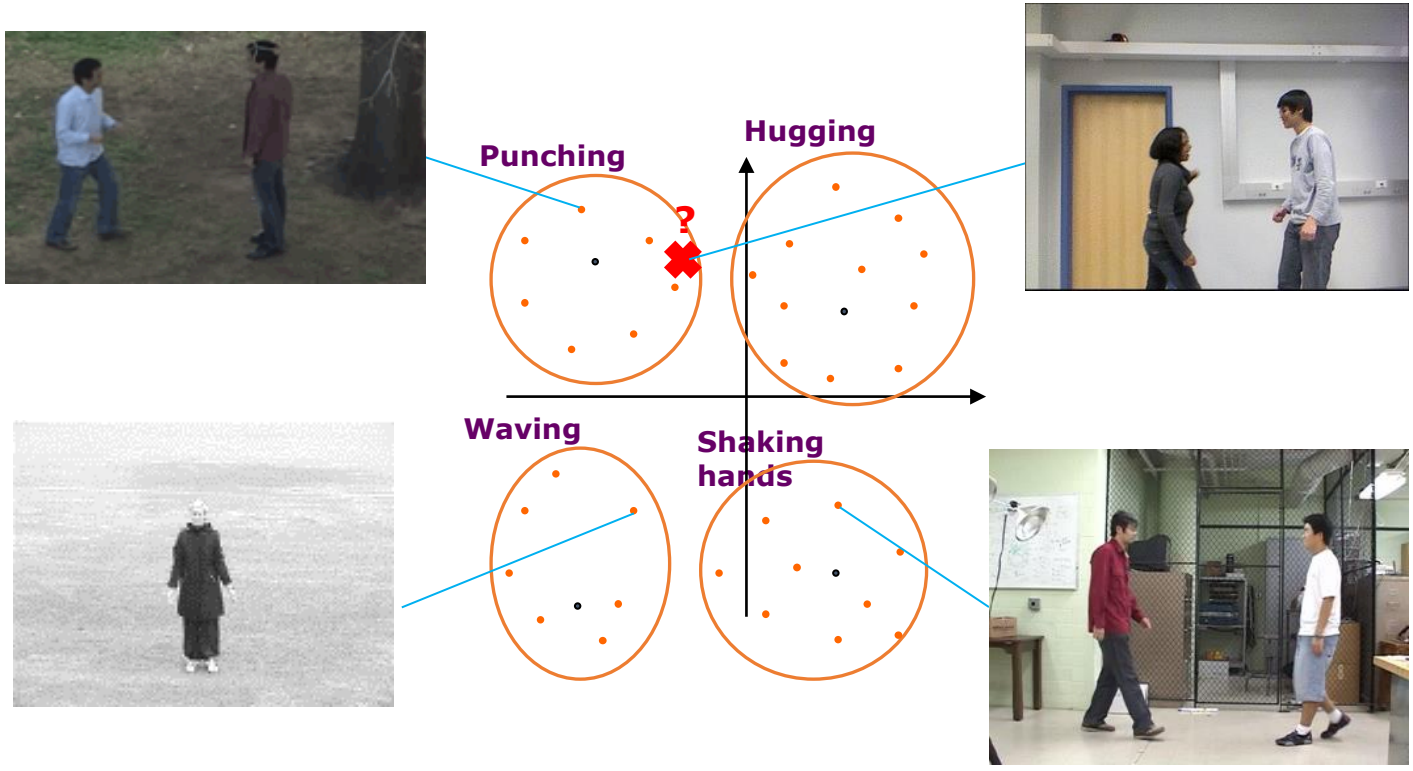

Example: Action Recognition

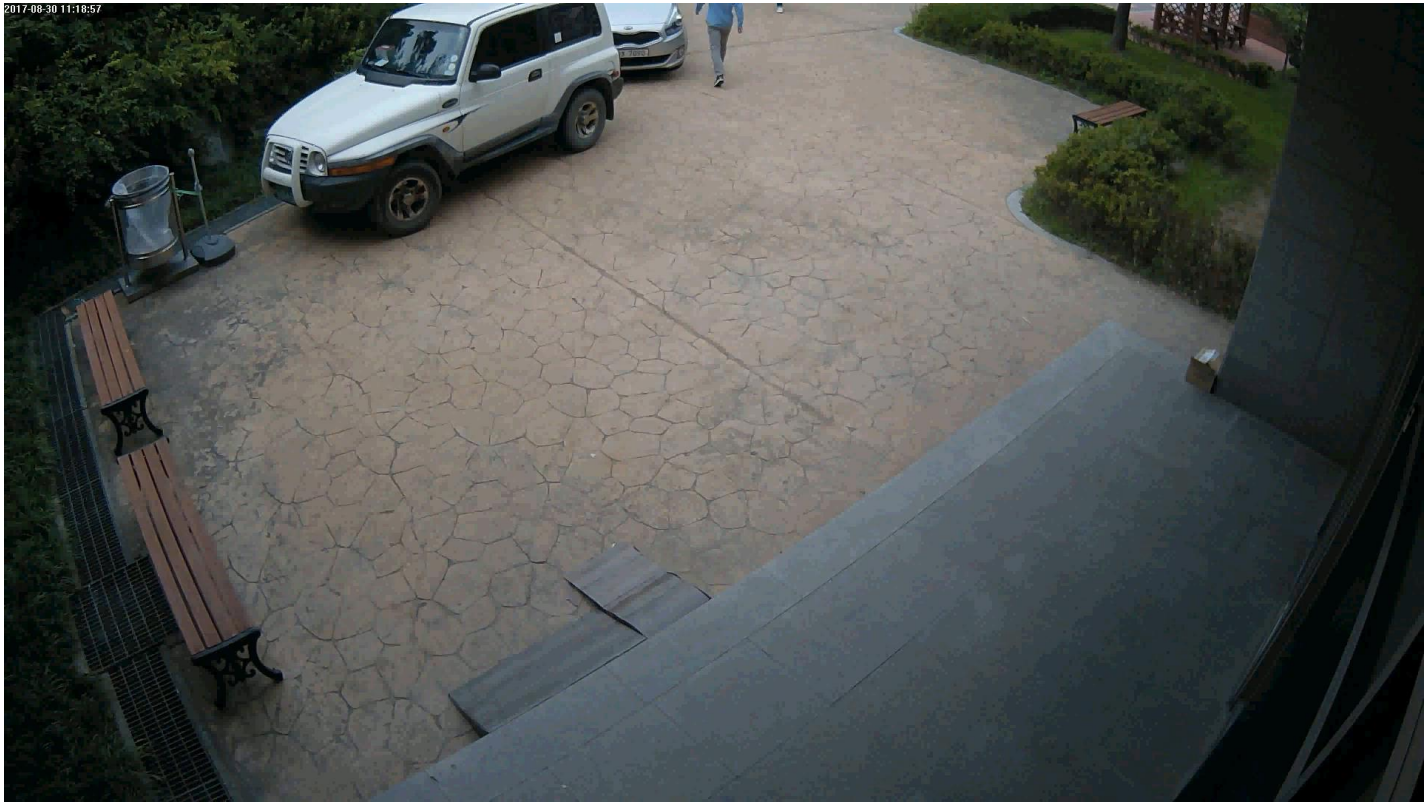

Group activity recognition

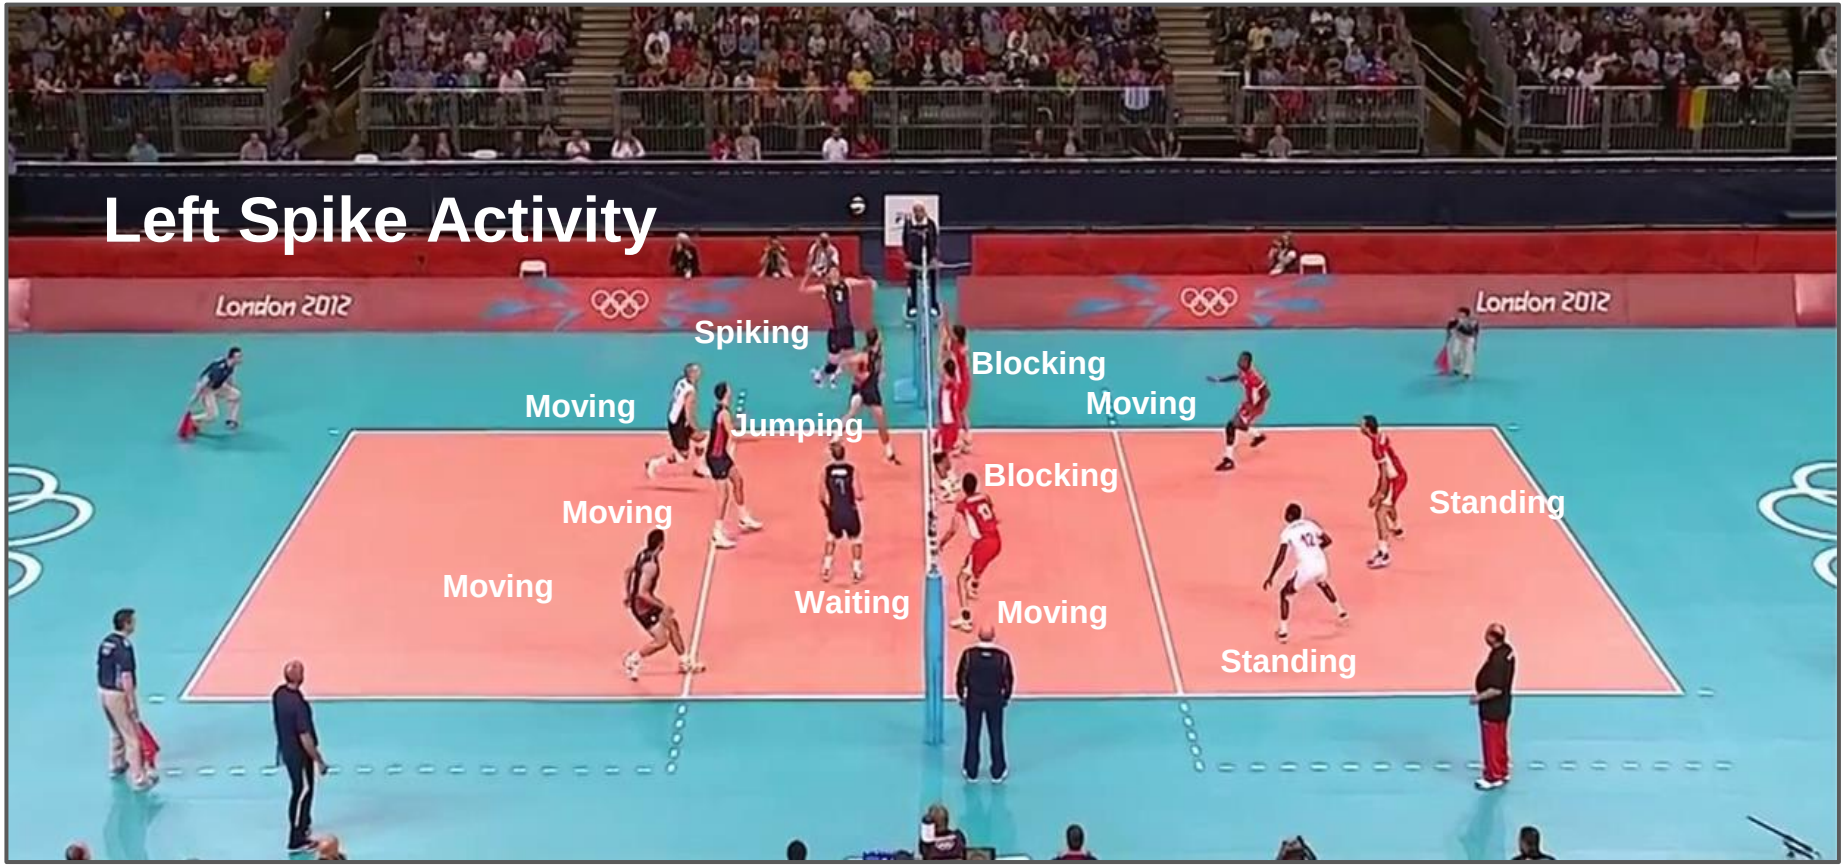

Relationships?

Left team:

- **Spike** by a player
- **Fake jump** by another player

Right team:

- 2 players jump to **block**

Learning Vision: Human vs. Computer

Computer Vision: Static Images (or Videos)

- Directly learn to detect objects
- Learn to segment

Human children: Sequence of related images

- Learn unity of objects
- Learn object permanence
- Learn to solve the binding problem

[slide credit: Sepp Hochreiter]

Major limitation of deep learning

Not data efficient: Learning requires **millions** of labeled examples, models do not generalize well to new domains; not like humans!

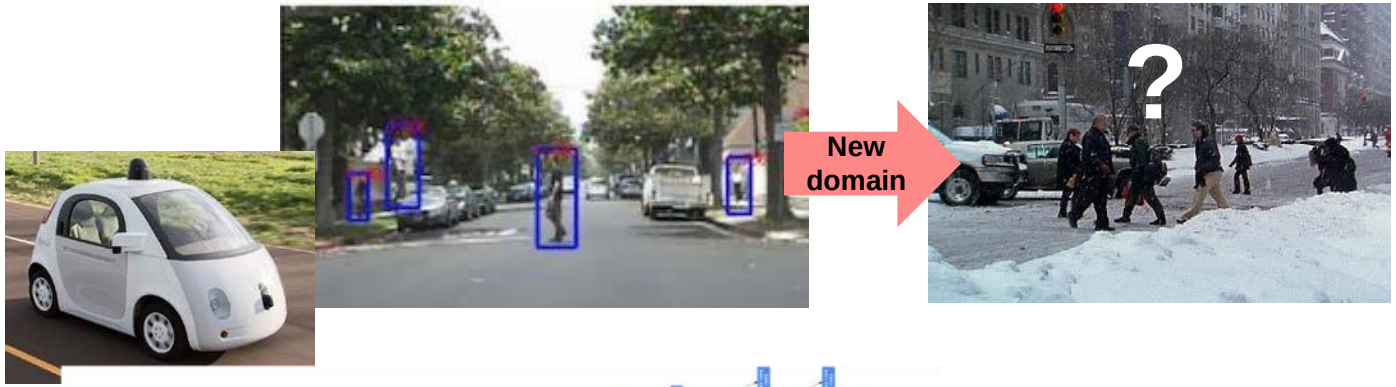

“What you saw is not what you get”

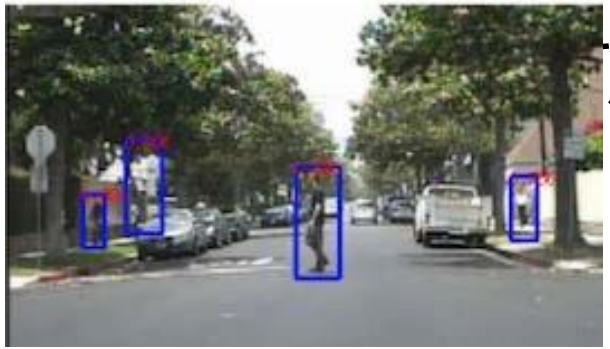

What your net is trained on

What it is asked to label

“Dataset Bias”
“Domain Shift”
“Domain Adaptation”
“Domain Transfer”

Example: Scene segmentation

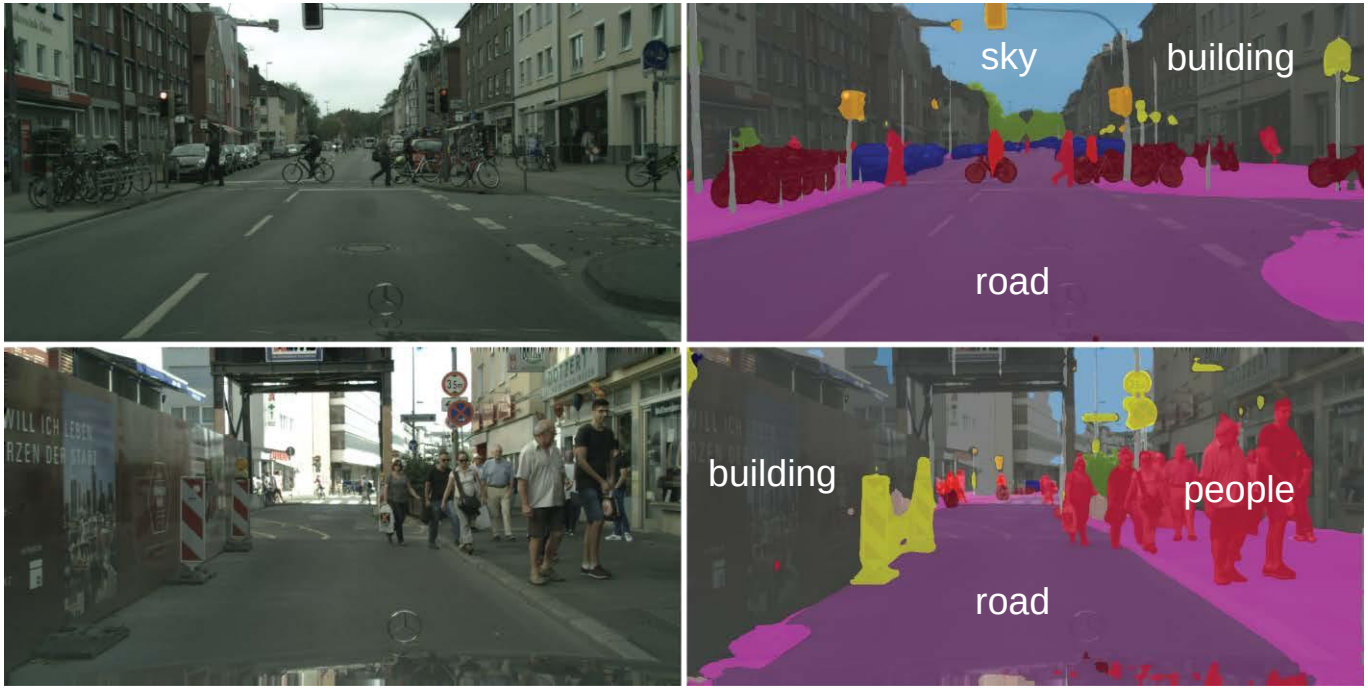

Train on Cityscapes, Test on Cityscapes

J. Hoffman, D. Wang, F. Yu, T. Darrell, FCNs in the Wild: Pixel-level Adversarial and Constraint-based Adaptation, Arxiv 2016

Domain shift: Cityscapes to SF

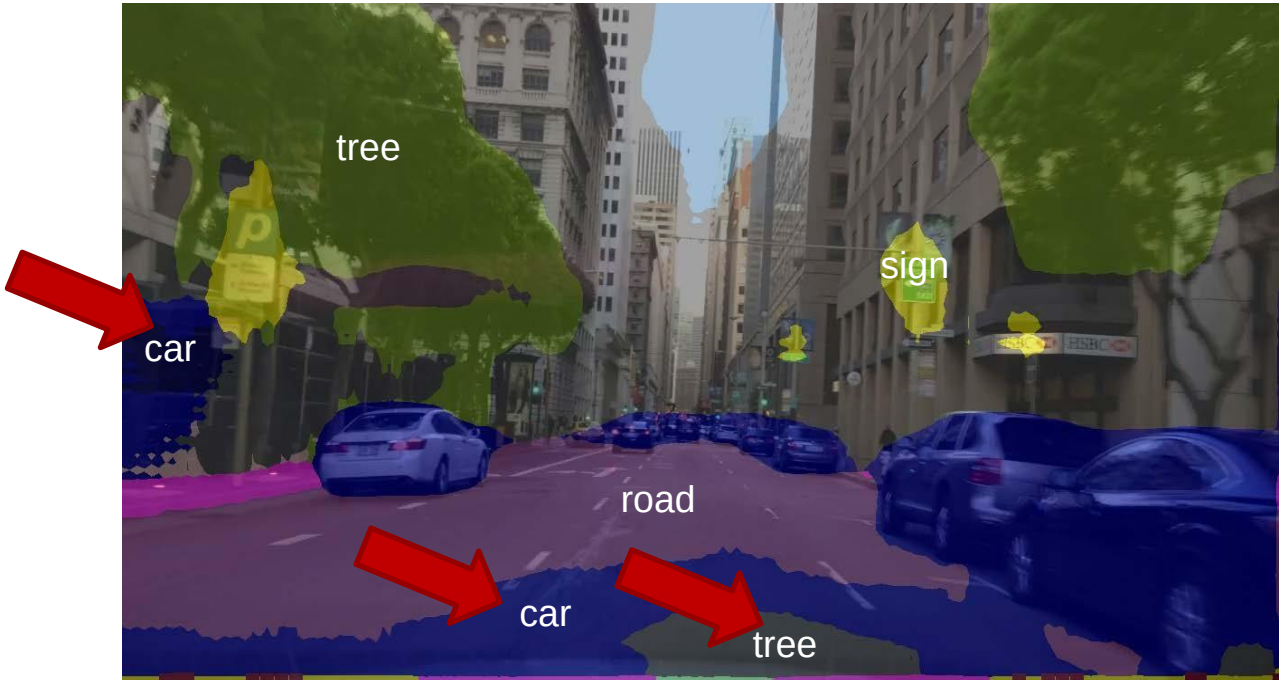

Train on Cityscapes, Test on San Francisco Dashcam

J. Hoffman, D. Wang, F. Yu, T. Darrell, FCNs in the Wild: Pixel-level Adversarial and Constraint-based Adaptation, Arxiv 2016

No tunnels in CityScapes?...

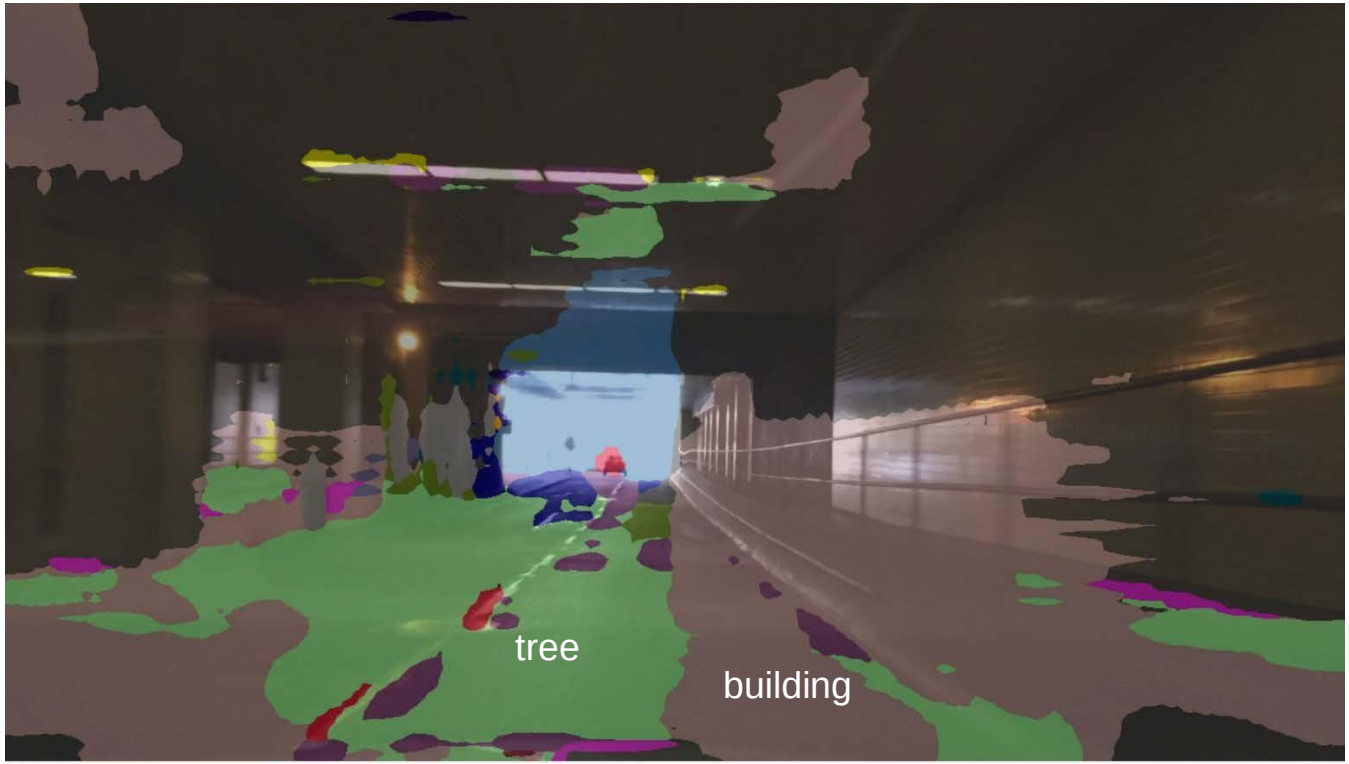

J. Hoffman, D. Wang, F. Yu, T. Darrell, FCNs in the Wild: Pixel-level Adversarial and Constraint-based Adaptation, Arxiv 2016

Domain Adaptation from source to target distribution

Adapt

backpack chair bike

? ? bike

Source Domain $\sim P_S(X, Y)$

lots of **labeled** data

$\neq$

Target Domain $\sim P_T(Z, H)$

unlabeled or limited labels

$$D_S = \{(\mathbf{x}_i, y_i), \forall i \in \{1, \ldots, N\}\}$$

$$D_T = \{(\mathbf{z}_j, ?), \forall j \in \{1, \ldots, M\}\}$$

Domain Shift: Image and Video Generation

Deep Generative Models for Image/Video Generation & Animation

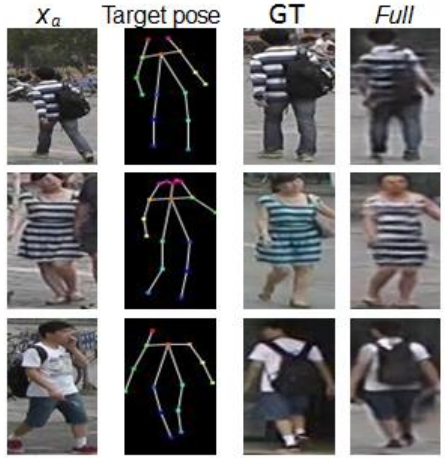

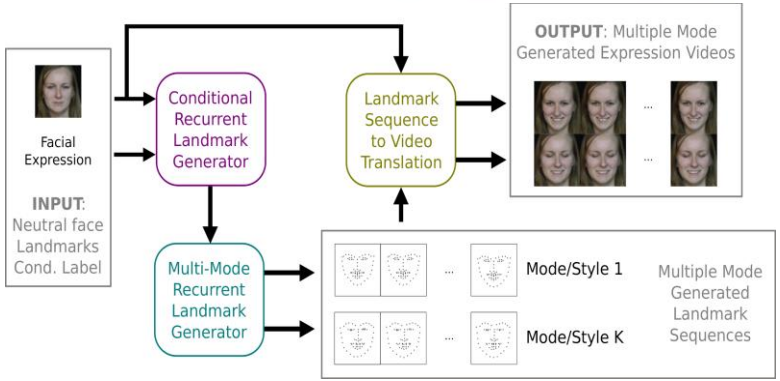

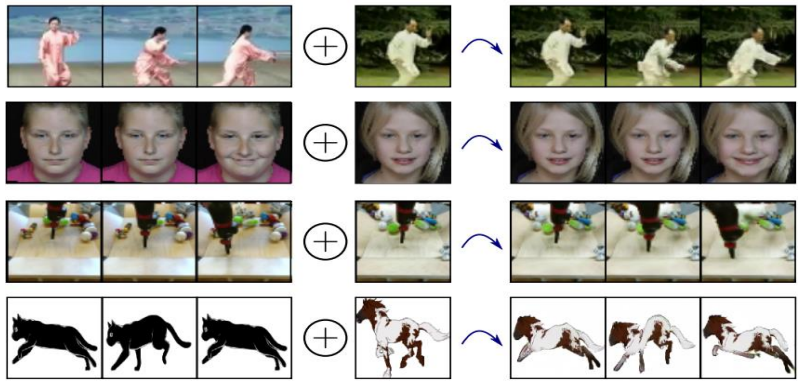

Arbitrary Object Animation without 3D modeling

Learned Keypoints

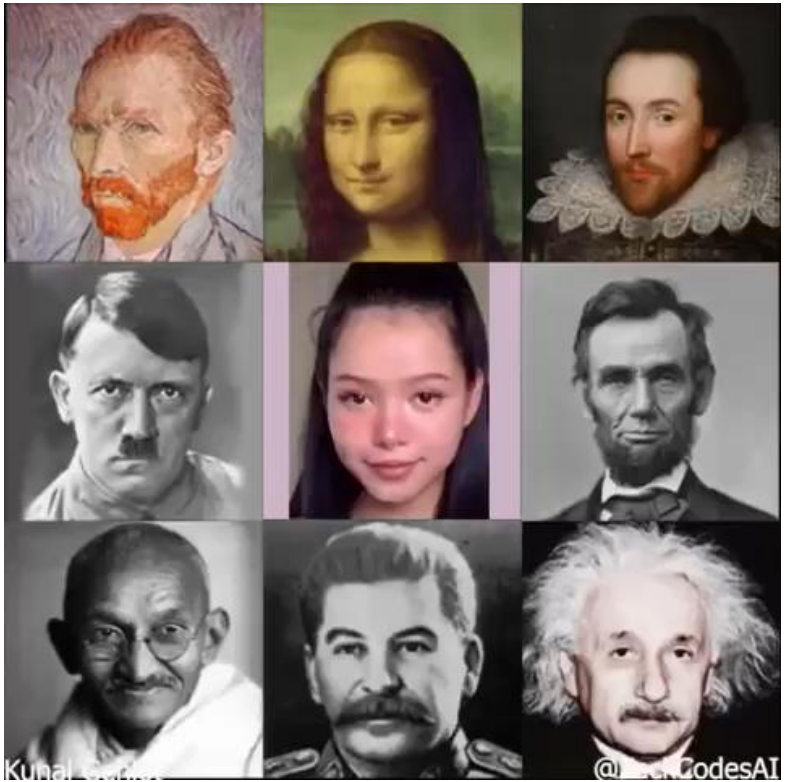
@CodesAI

Deep Fakes …


BANCA NAȚIONALĂ A ROMÂNIEI
5
GUVERNATOR
CASIER CENTRAL
5 LEI CINCI
CINCI LEI

Tai-Chi-HD

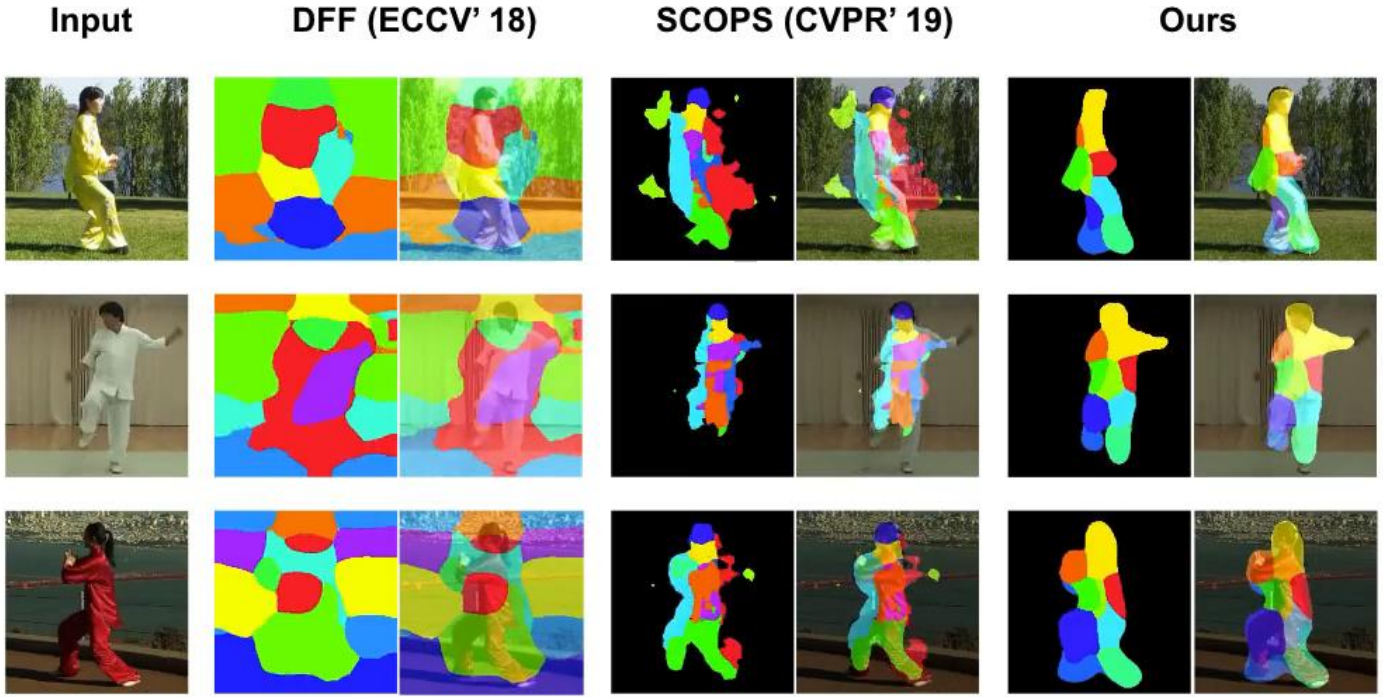


humans
SCALE
talkens.ai
TALKENS.AI

Playable Video Generation

Menapace, et al., “Playable Video Generation”, in CVPR 2021

https://github.com/willi-menapace/PlayableVideoGeneration

Playable Environments

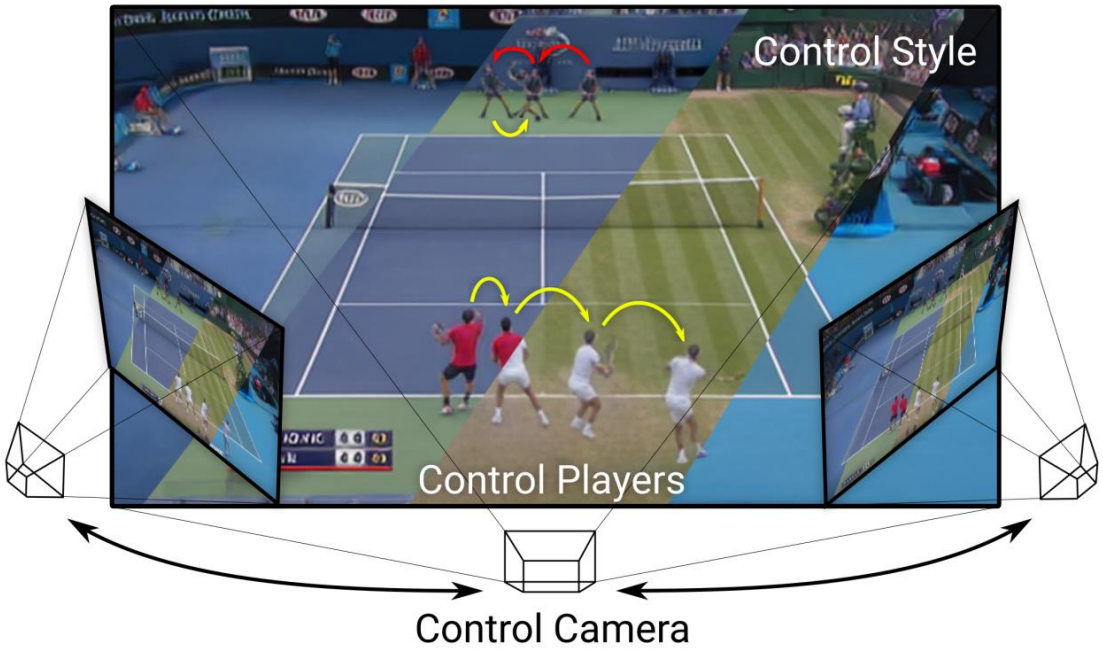

- Learn a model that represents the observed environment
- Allow the user to input actions to the model through a controller at test time

Playable Environments

Menapace et al, “Playable Environments: Video Manipulation in Space and Time”, in CVPR 2022

https://github.com/willi-menapace/PlayableEnvironments

Playable Environments

Menapace et al, “Playable Environments: Video Manipulation in Space and Time”, in CVPR 2022

https://github.com/willi-menapace/PlayableEnvironments

Playable Environments

Menapace et al, “Playable Environments: Video Manipulation in Space and Time”, in CVPR 2022

https://github.com/willi-menapace/PlayableEnvironments

Topics

- Scene recognition and understanding
- Object detection/localization/recognition
- Pixel-level prediction
  - Semantic segmentation
  - Depth estimation
- Motion analysis and activity recognition
- Biometrics: face, gesture, body pose, emotions
- Social signals processing

Topics

- Image and video generation
- Tiny ML for computer vision
- Transfer learning and deep domain adaptation
- Explainable AI
- Privacy, transparency and ethics in vision

Course organization

- Alternate oral presentations with practical sessions:
  - We will present in detail some of the basic code implementation to provide details on how things are done in practice
- Final evaluation
  - Students (groups) will be asked to choose from a list of recent important articles
  - Briefly present the main ideas
  - Propose extensions that could improve/generalize the approach

Course timeline

- Classes on:
  - September 12, 13, 19, 20, 26, 27
  - October 3, 4, 17, 18 (pause for ACM MM and ECCV)
  - November 7, 8, 14, 15, 22, 28, 29
- Student presentations on:
  - December 6, 12, 13

Scene Understanding and Object Recognition

Why is Vision Interesting?

- Psychology
  - ~ 35% of cerebral cortex is for vision
  - Vision is how we experience the world
- Engineering
  - Want machines to interact with the world
  - Digital images are everywhere

Measurement vs. Perception

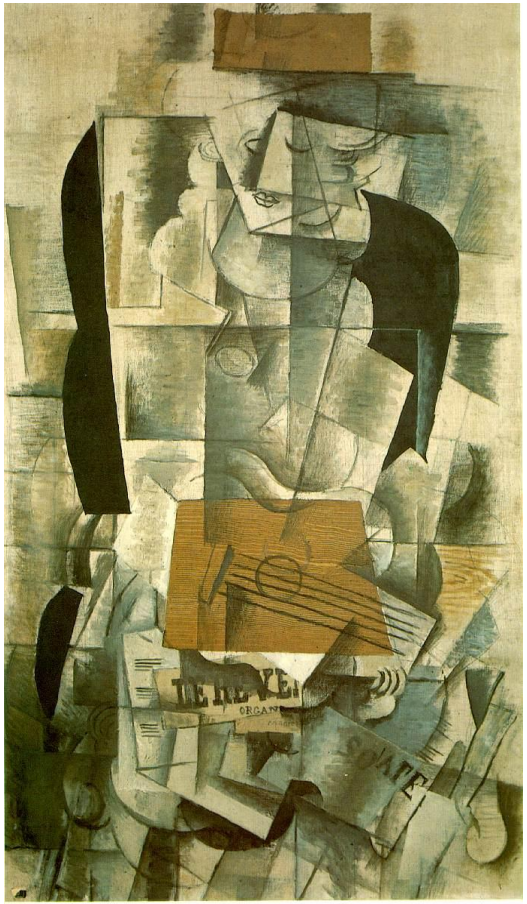

Brightness: Measurement vs. Perception

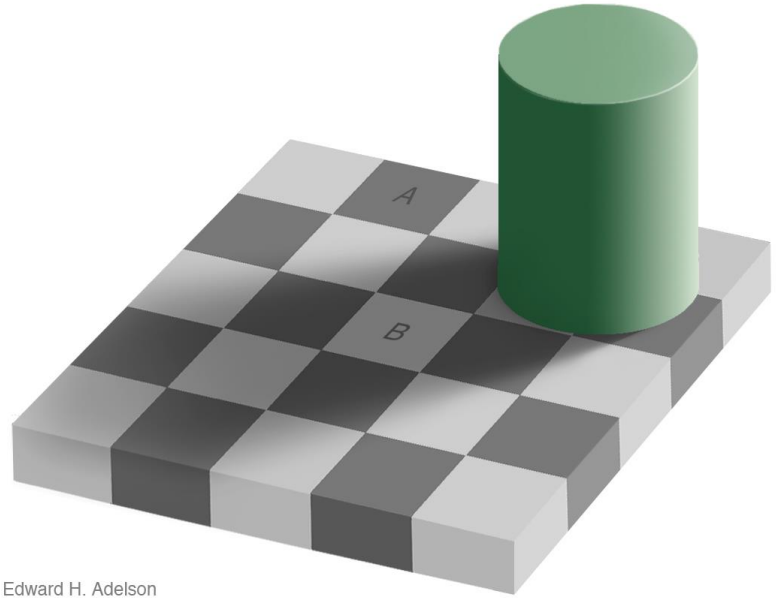

Edward H. Adelson

Brightness: Measurement vs. Perception

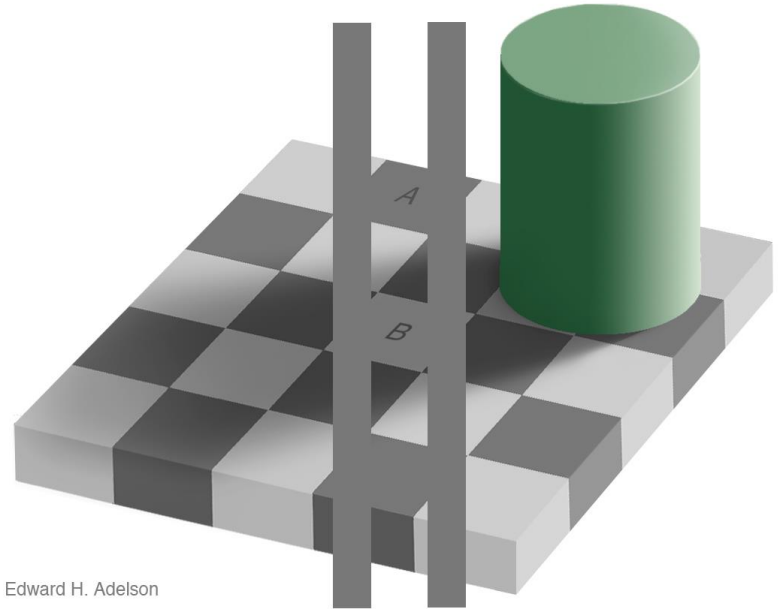

Proof!

Lengths: Measurement vs. Perception

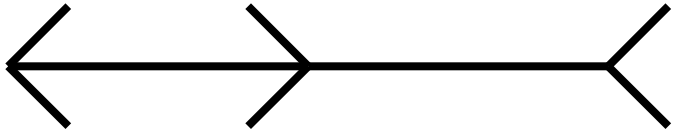

Müller-Lyer Illusion

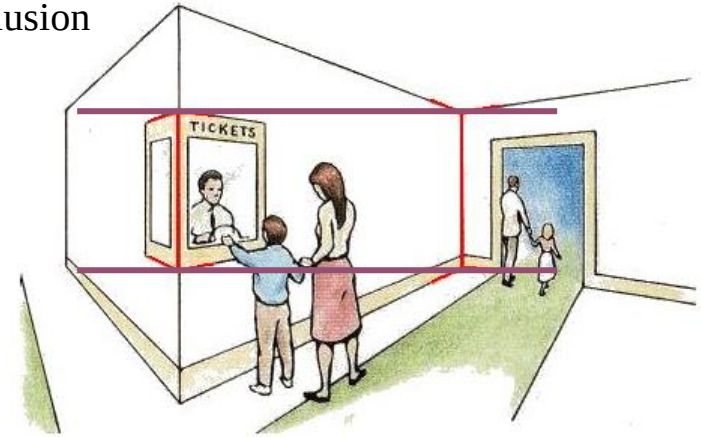

https://michaelbach.de/ot/sze-muelue/index.html

Vision is Inferential: Prior Knowledge

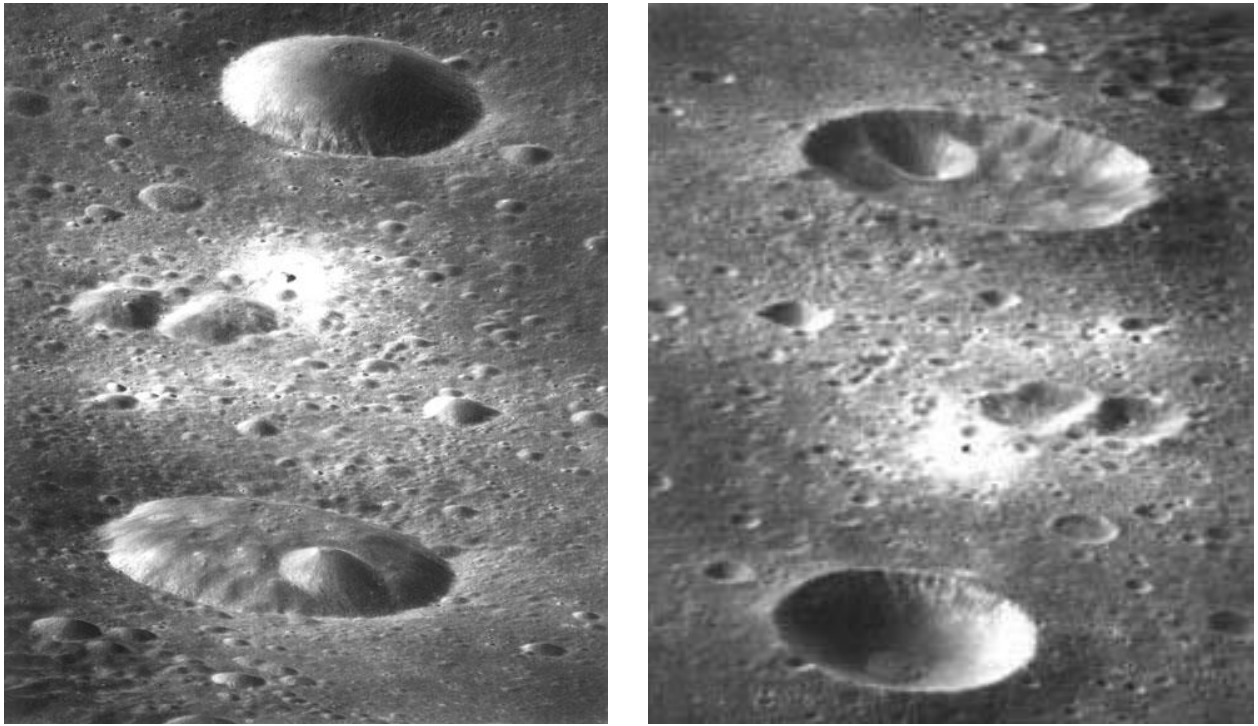

Vision is Inferential: Prior Knowledge

- Our visual system is very sophisticated
- Humans can interpret images successfully under a wide range of conditions – even in the presence of very limited cues

Playing with Perspective

- Perspective gives us very strong depth cues
  ⇨ hence we can perceive a 3D scene by viewing its 2D representation (i.e. image)
- An example where perception of 3D scenes is misleading:

“Ames room”

A clip from "The computer that ate Hollywood" documentary. Dr. Vilayanur S. Ramachandran.

- *“What if I don’t care about this wishy-washy human perception stuff? I just want to make my robot go!”*
- Small Reason:
  - For measurement, other sensors are often better (in DARPA Grand Challenge, vision was barely used!)
  - For navigation, you still need to learn!
- Big Reason:
- The goals of computer vision (**what + where**) are in terms of what humans care about

So what do humans care about?

Verification: is that a bus?

Detection: are there cars?

Identification: is that a picture of Mao?

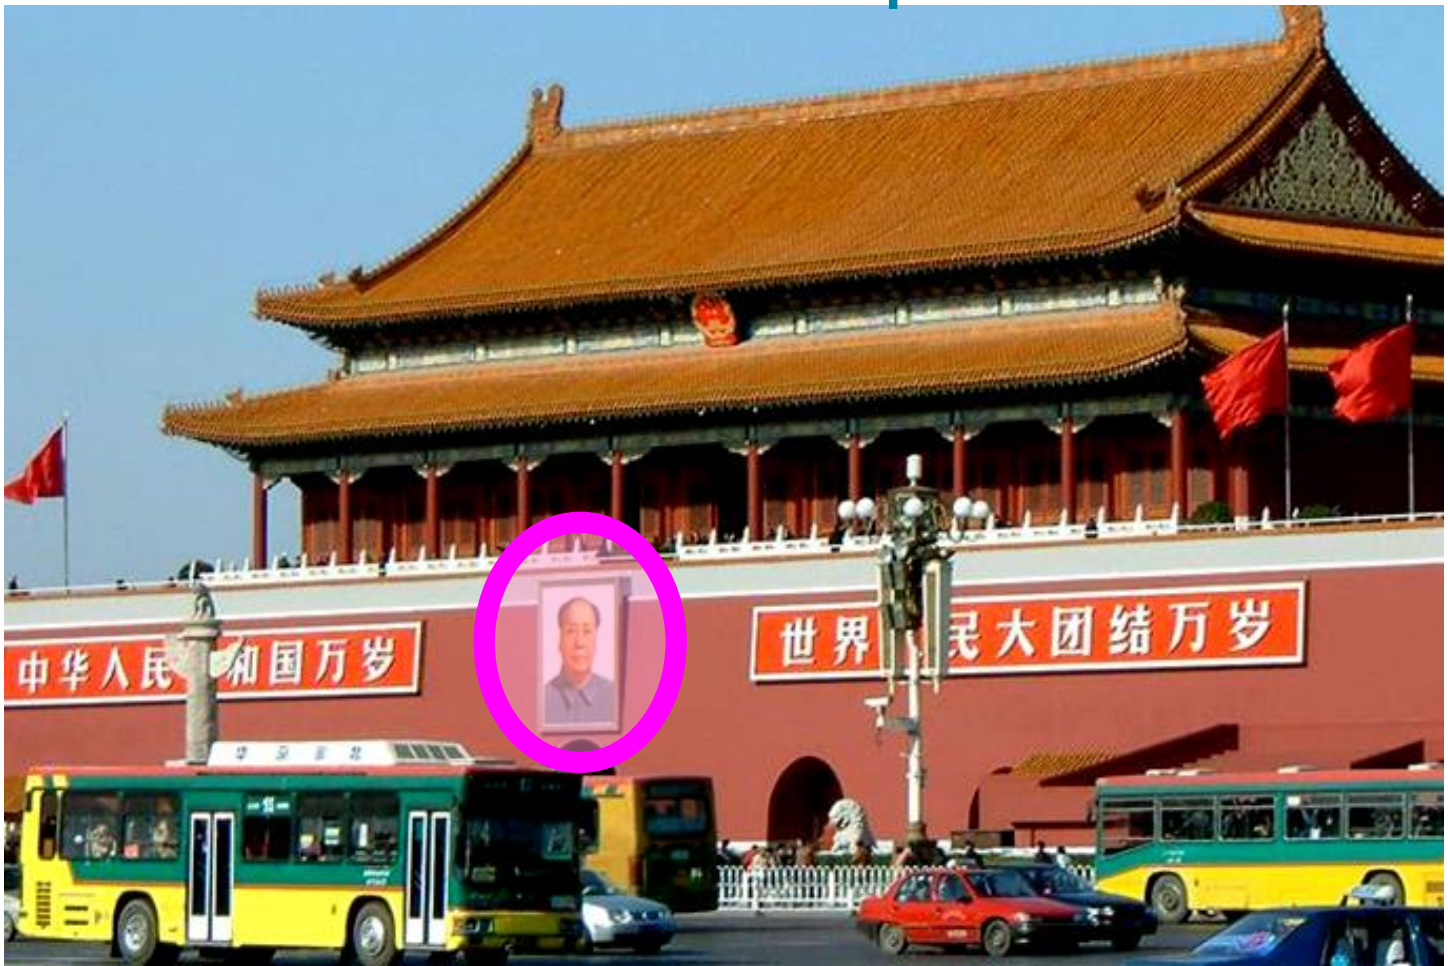

Object categorization

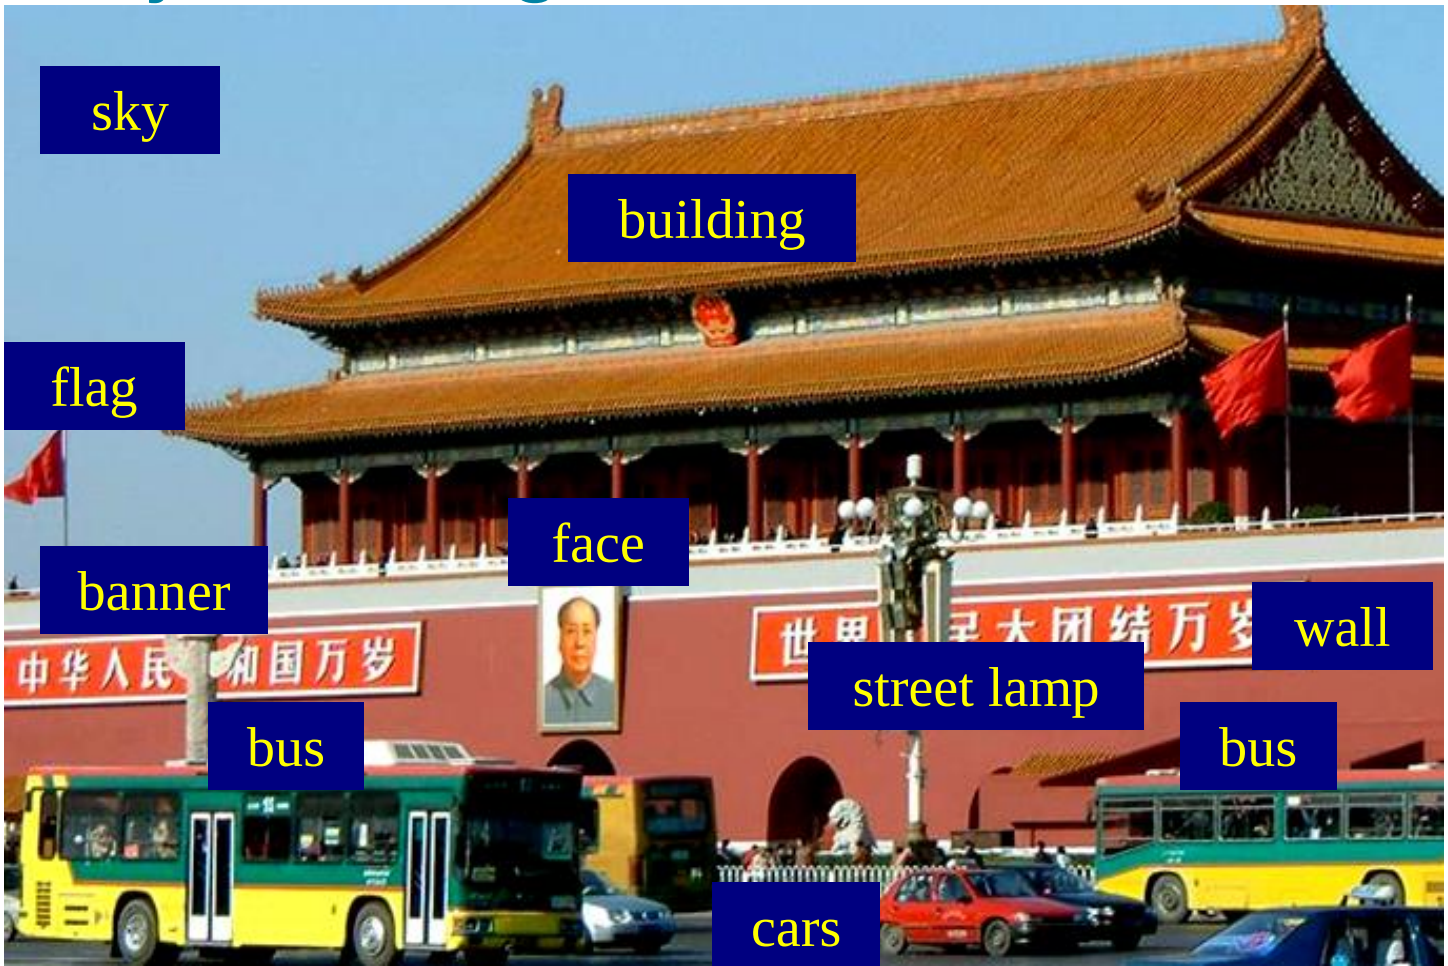

Scene and context categorization

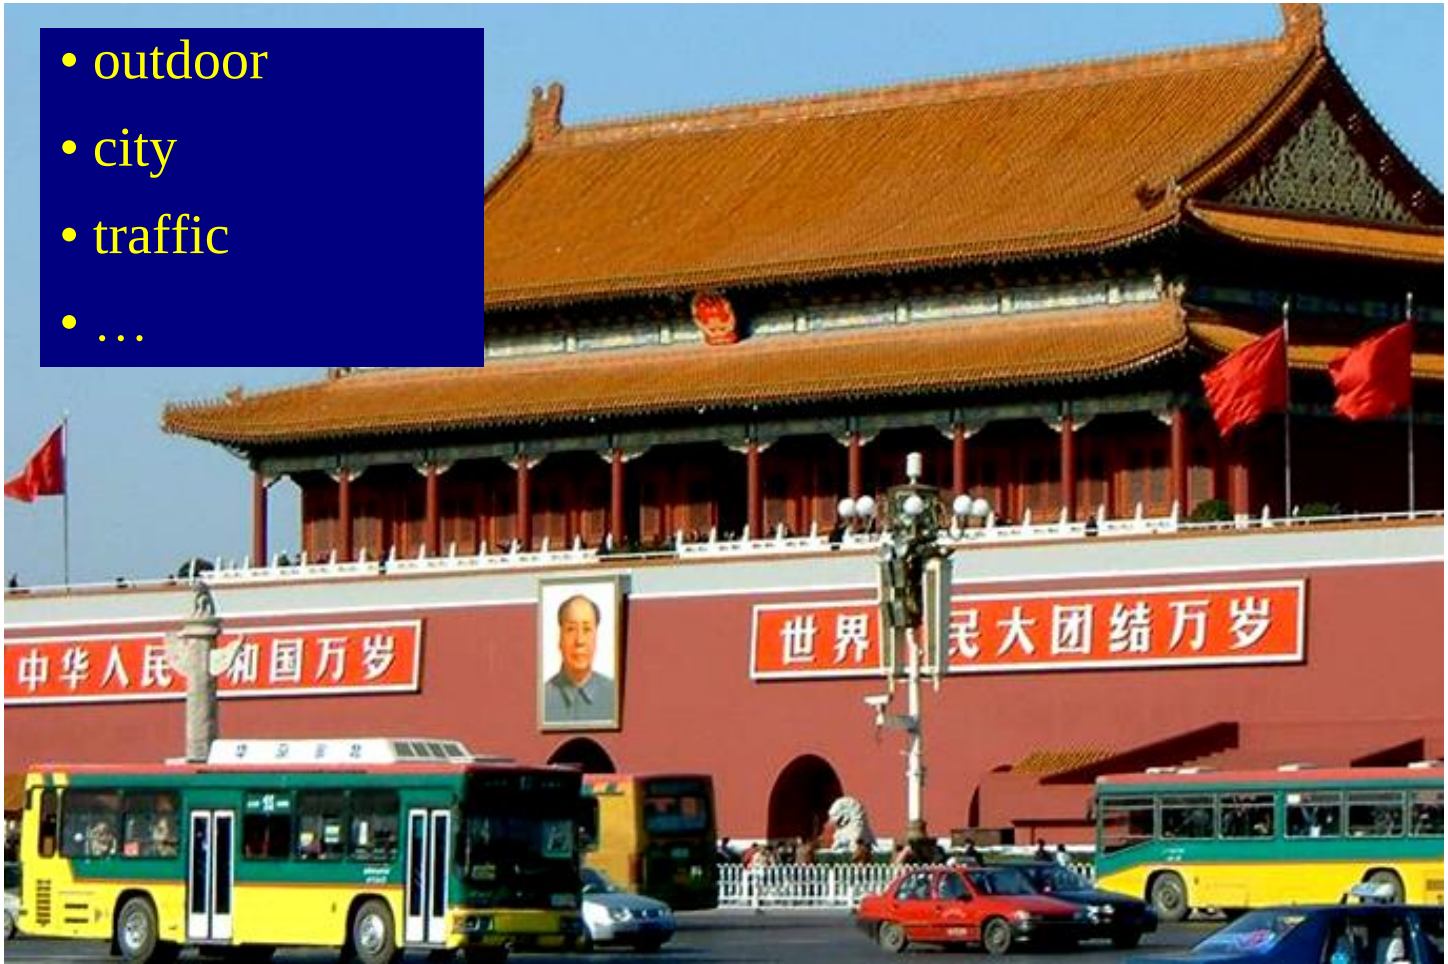

Rough 3D layout, depth ordering

Challenges 1: view point variation

Michelangelo 1475-1564

Challenges 2: illumination

Challenges 3: occlusion

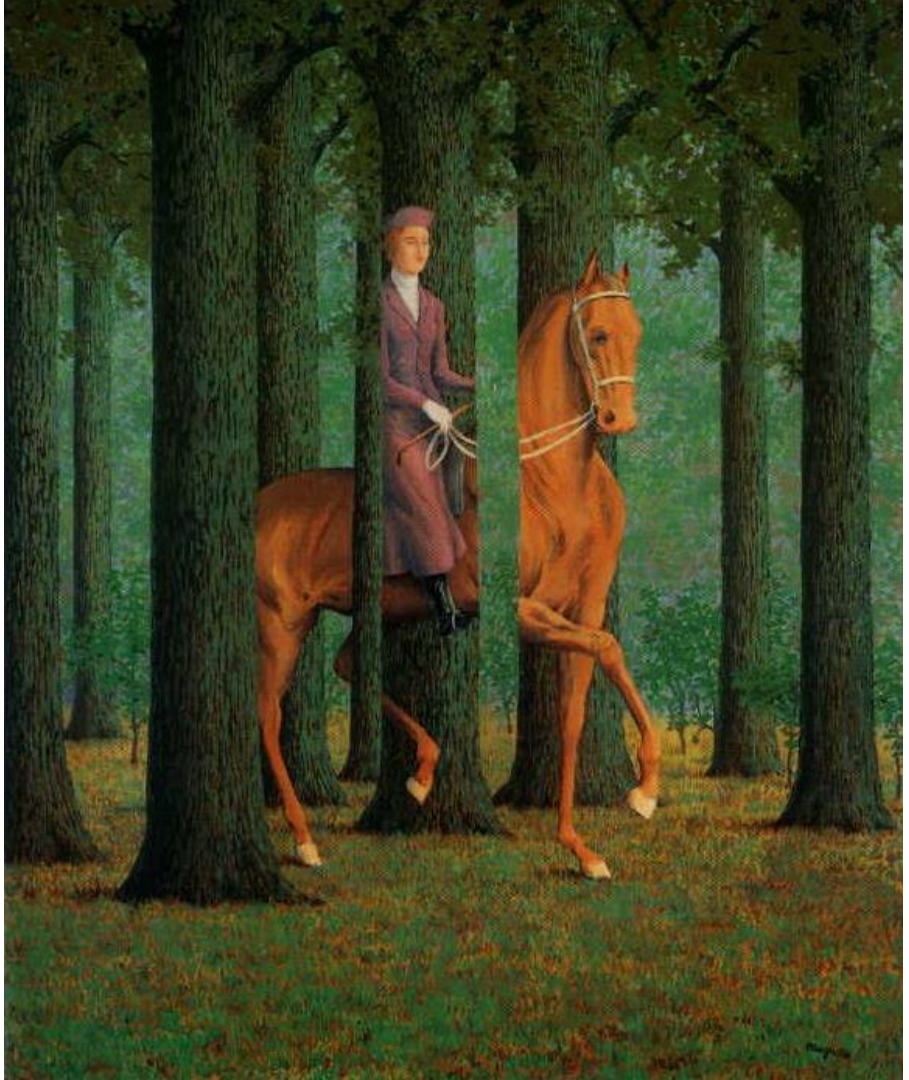

Magritte, 1957

Challenges 4: scale

Challenges 5: deformation

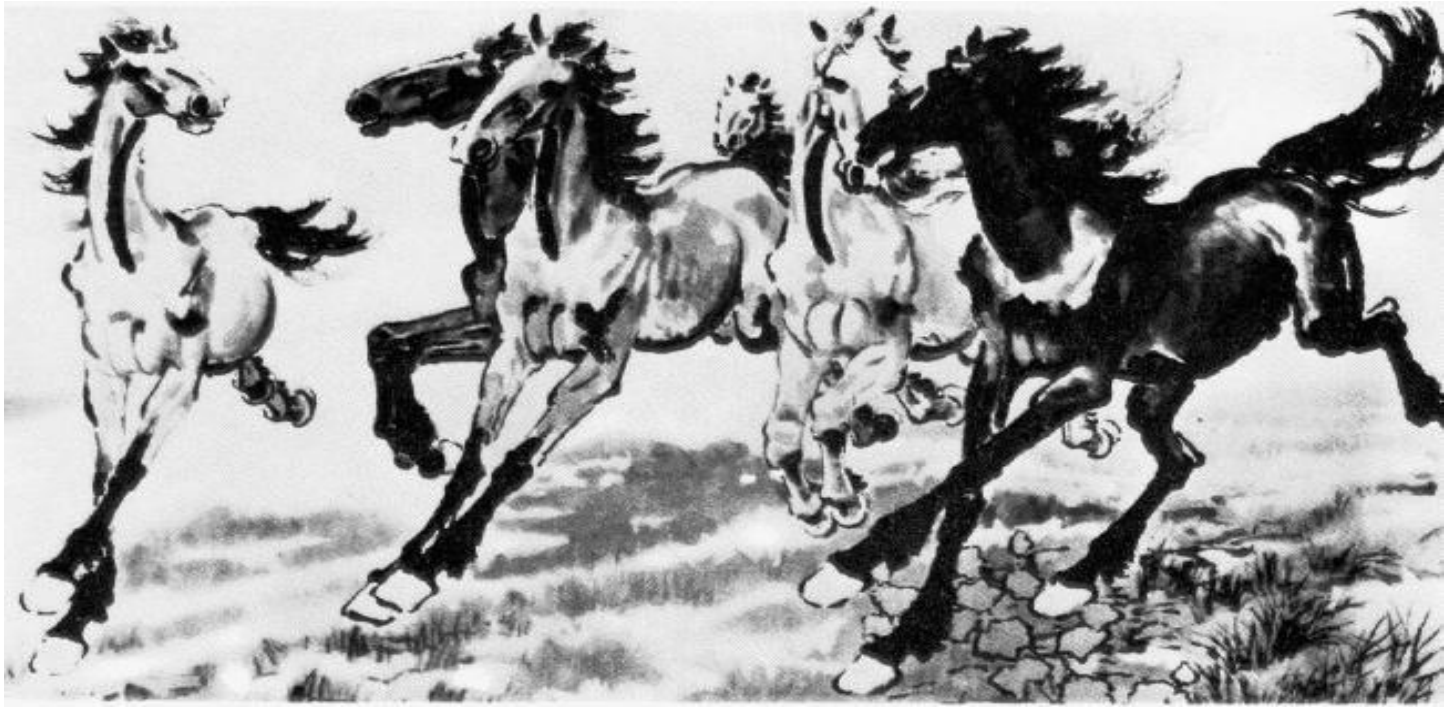

Xu, Beihong 1943

Challenges 6: background clutter

Klimt, 1913

Challenges 7: object intra-class variation

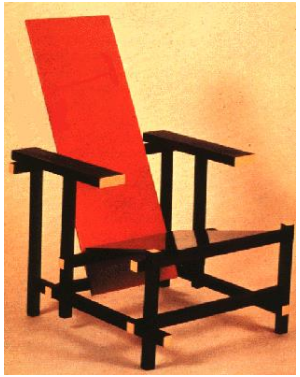

Challenges 8: local ambiguity

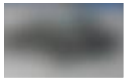

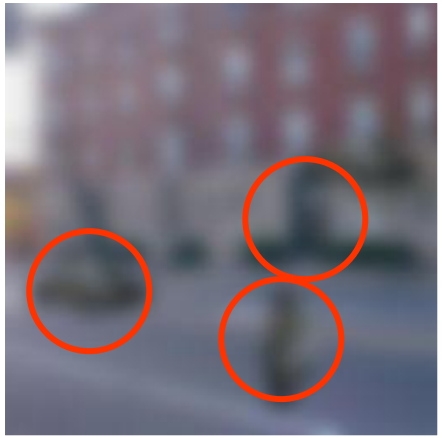

Object detection and localization

Computer Vision Tasks

Classification

CAT

Classification + Localization

CAT

Object Detection

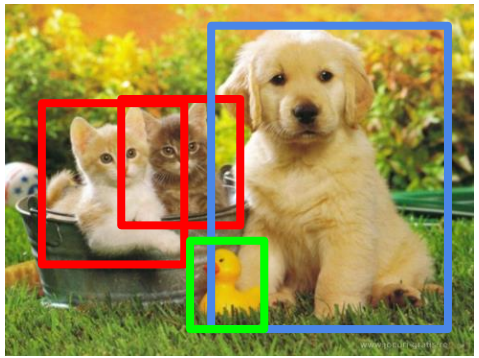

CAT, DOG, DUCK

Instance Segmentation

CAT, DOG, DUCK

Single object

Multiple objects

Faces

- Face detection
- One category: face
- Frontal faces
- Fairly rigid, unoccluded

1990’s

Human Face Detection in Visual Scenes. H. Rowley, S. Baluja, T. Kanade. 1995.

Faces

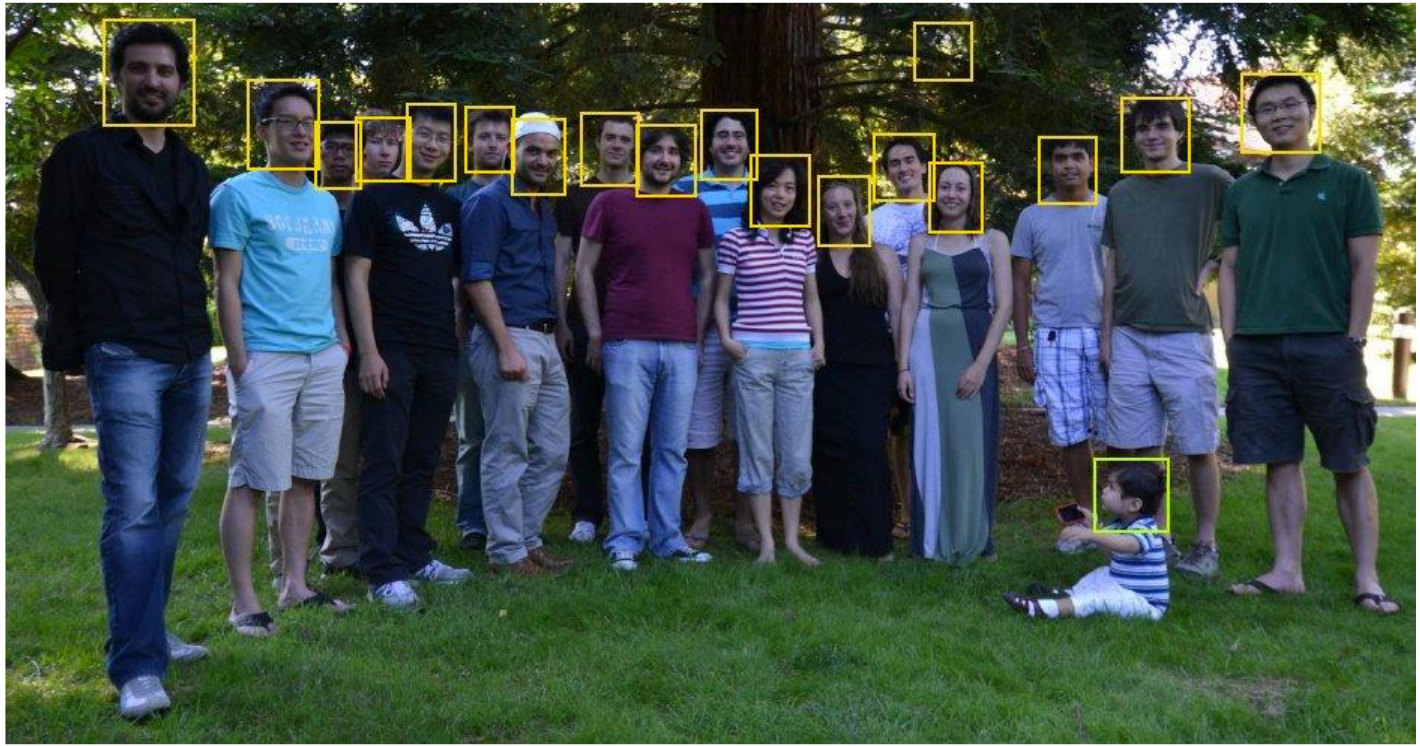

Pedestrians

- One category: pedestrians
- Slight pose variations and small distortions
- Partial occlusions

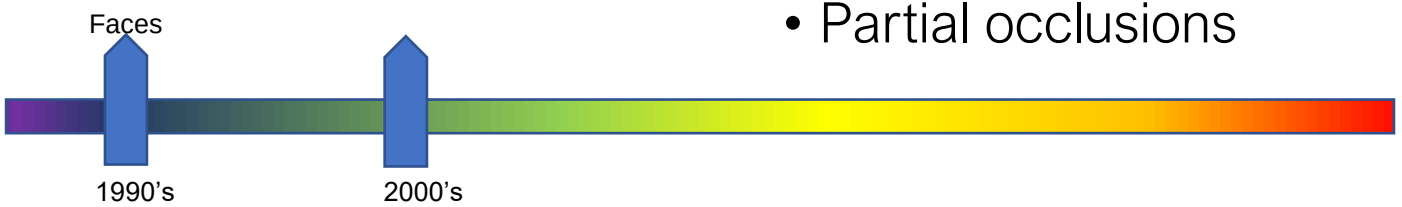

Histograms of Oriented Gradients for Human Detection. N. Dalal and B. Triggs. CVPR 2005

Pedestrians

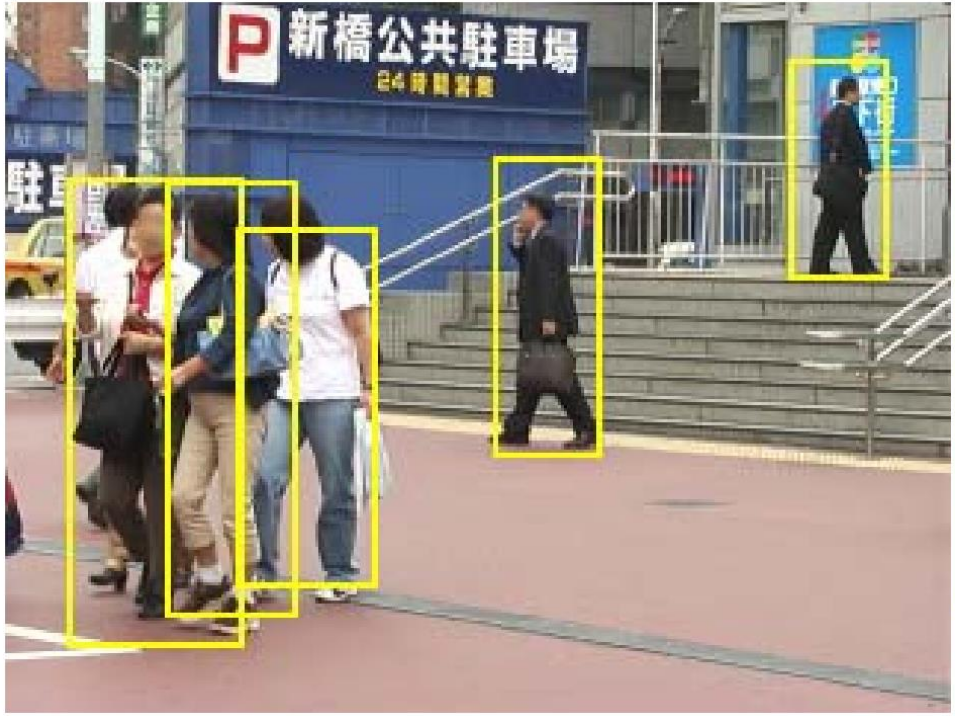

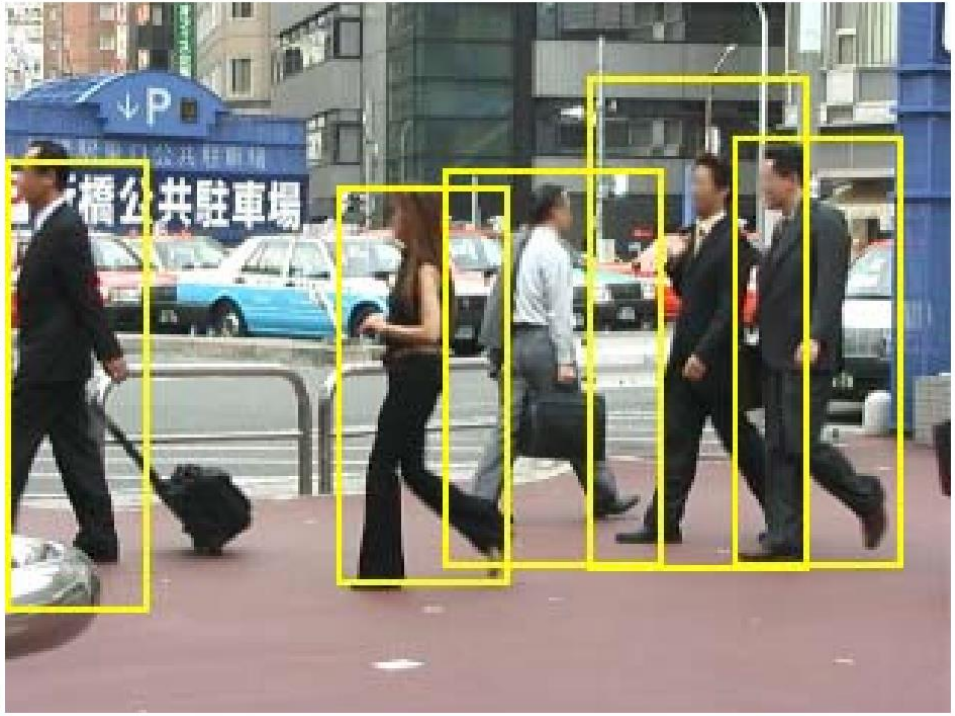

Applications: Tagging People

Applications: Autonomous Driving

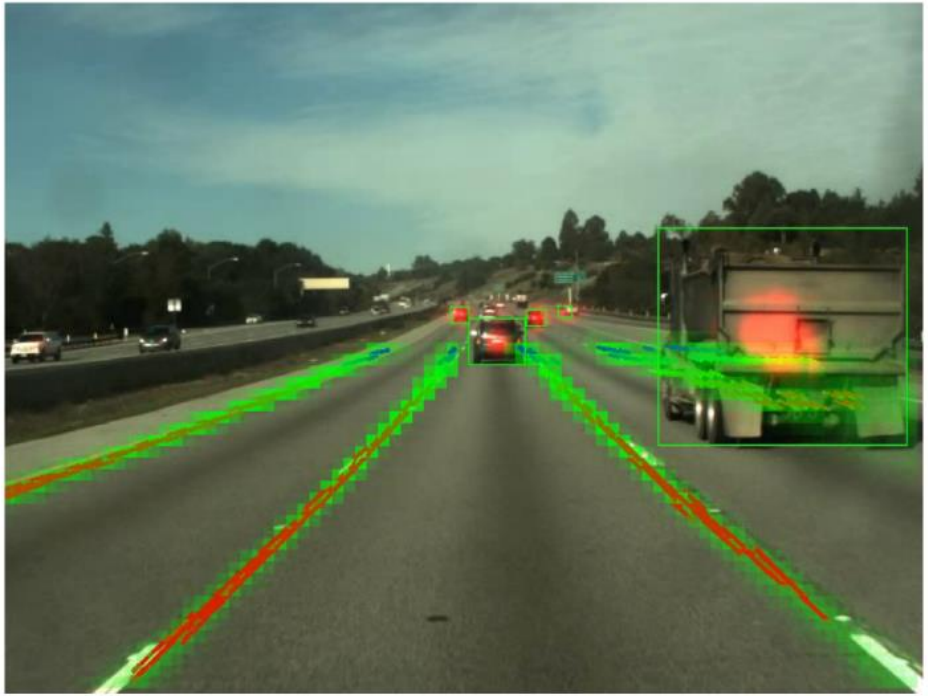

Applications: Robotics

Applications: Tracking

Applications: Semantic Segmentation

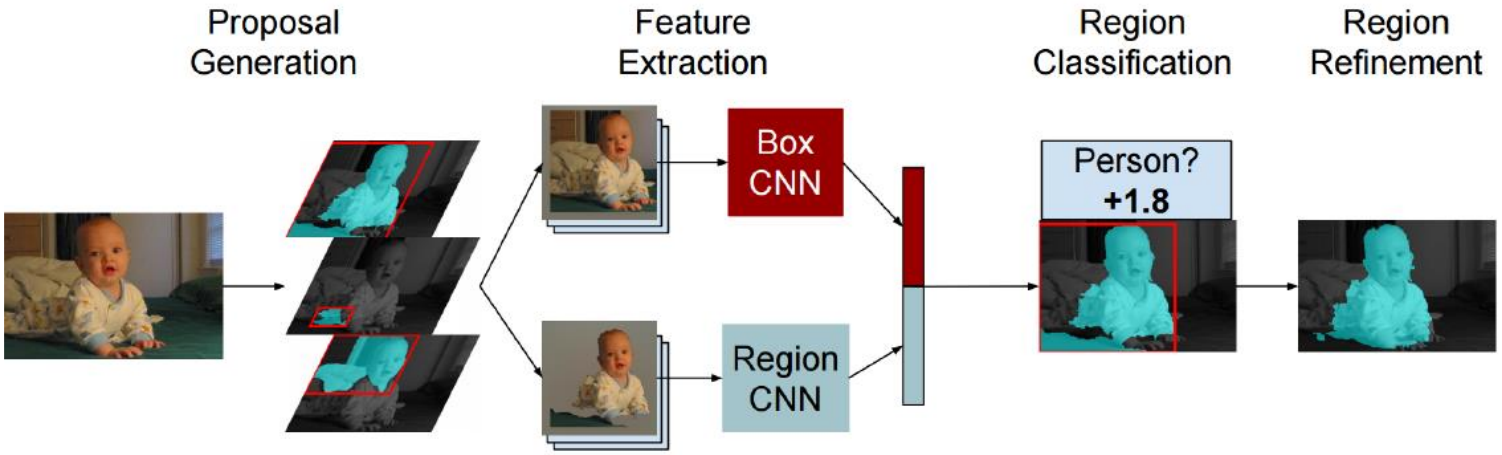

PASCAL VOC

- 20 categories
- 10K images
- Large pose variations, heavy occlusions
- Generic scenes
- Cleaned up performance metric

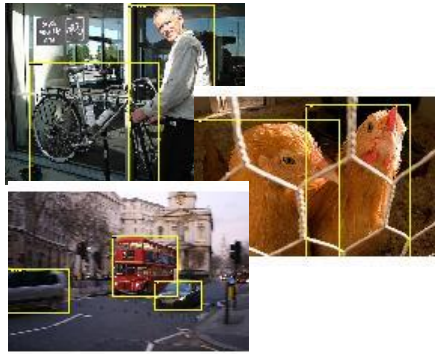

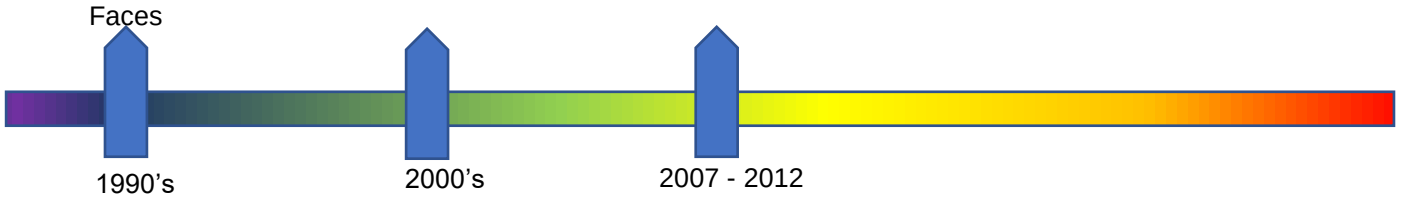

Coco

- 80 diverse categories
- 100K images
- Heavy occlusions, many objects per image, large scale variations

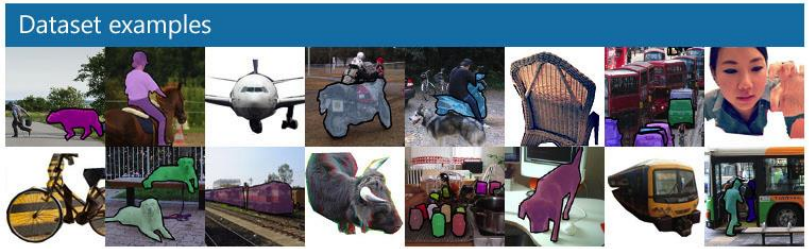

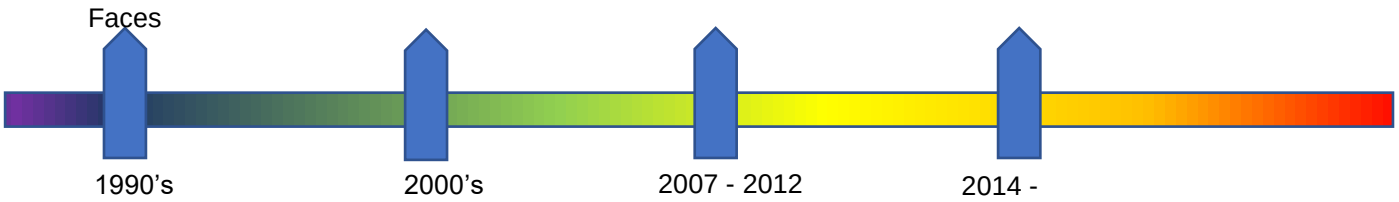

Evaluation metric

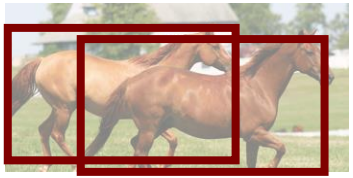

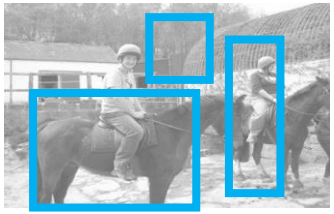

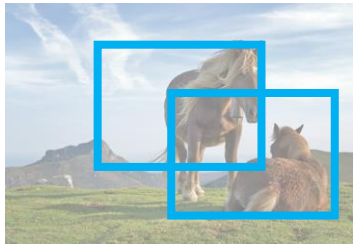

Matching detections to ground truth

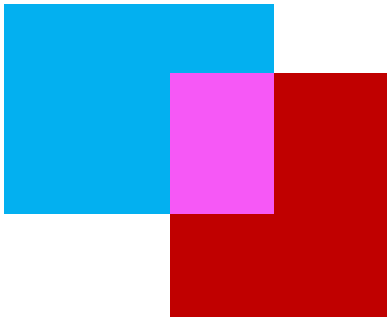

$$IoU(A, B) = \frac{|A \cap B|}{|A \cup B|}$$

Matching detections to ground truth

- Match detection to most similar ground truth
  - highest IoU
- If IoU > 50%, mark as correct
- If multiple detections map to same ground truth, mark only one as correct
- **Precision** = #correct detections / total detections
- **Recall** = #ground truth with matched detections / total ground truth

Tradeoff between precision and recall

- ML usually gives scores or probabilities, so we need to threshold
- Too low threshold → too many detections → low precision, high recall
- Too high threshold → too few detections → high precision, low recall
- Right tradeoff depends on application
  - Detecting cancer cells in tissue: need high recall
  - Detecting edible mushrooms in forest: need high precision

Average precision

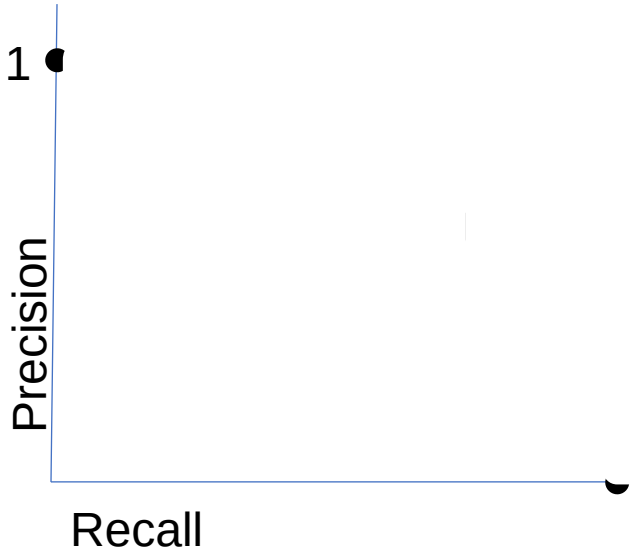

Average precision

Average average precision

- AP marks detections with overlap > 50% as correct
- But may need better localization
- *Average* AP across multiple overlap thresholds
- Confusingly, still called average precision
- Introduced in COCO

Mean and category-wise AP

- Every category evaluated independently
- Typically report mean AP averaged over all categories
- Confusingly called “mean Average Precision”, or “mAP”

Why is detection hard(er)?

- Precise localization

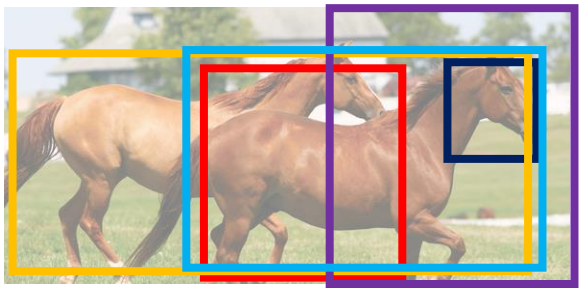

Why is detection hard(er)?

- Much larger impact of pose

Why is detection hard(er)?

- Occlusion makes localization difficult

Why is detection hard(er)?

- Counting

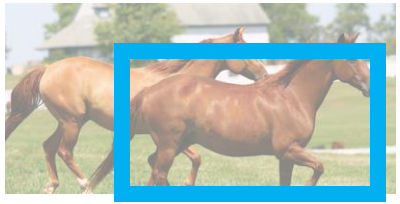

Why is detection hard(er)?

- Small objects