

# Object detection and localization

# Computer Vision Tasks

## Classification




CAT

## Classification + Localization



CAT

## Object Detection



CAT, DOG, DUCK

## Instance Segmentation



CAT, DOG, DUCK

Single object

Multiple objects

# Faces



- Face detection
- One category: face
- Frontal faces
- Fairly rigid, unoccluded



Human Face Detection in Visual Scenes. H. Rowley, S. Baluja, T. Kanade. 1995.


# Faces



# Pedestrians





- One category: pedestrians
- Slight pose variations and small distortions
- Partial occlusions



Histograms of Oriented Gradients for Human Detection. N. Dalal and B. Triggs. CVPR 2005




# Pedestrians



# Applications: Tagging People




# Applications: Autonomous Driving







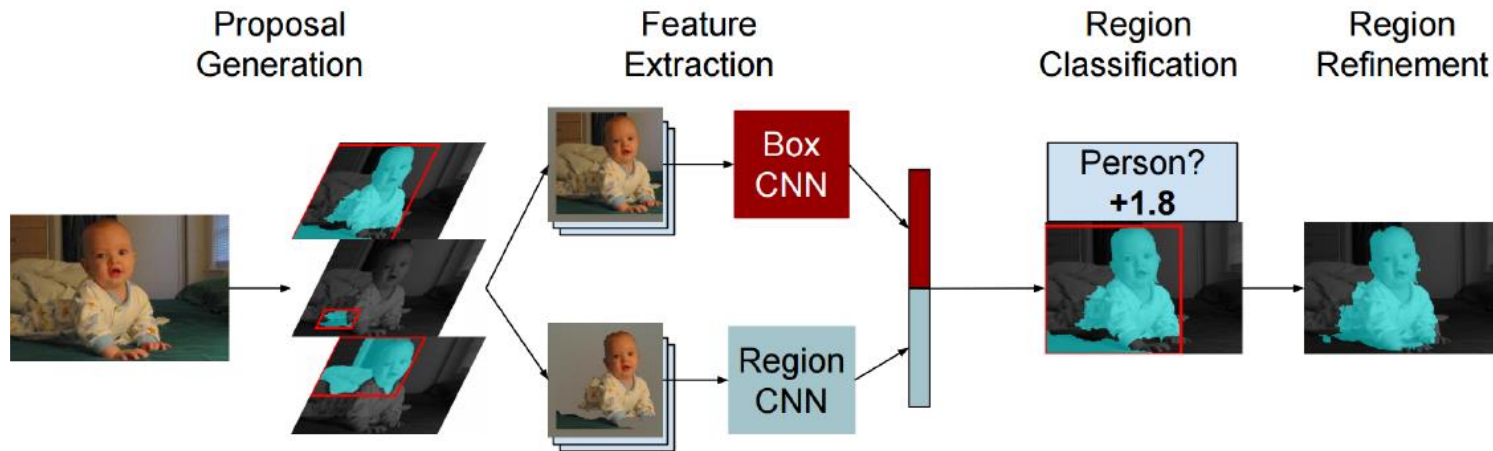
# Applications: Robotics



# Applications: Tracking





# Applications: Semantic Segmentation





# PASCAL VOC

- 20 categories
- 10K images
- Large pose variations, heavy occlusions
- Generic scenes
- Cleaned up performance metric







# Coco

- 80 diverse categories
- 100K images
- Heavy occlusions, many objects per image, large scale variations






# Evaluation metric



# Matching detections to ground truth



$$IoU(A, B) = \frac{|A \cap B|}{|A \cup B|}$$


# Matching detections to ground truth

- Match detection to most similar ground truth
  - highest IoU
- If  $\text{IoU} > 50\%$ , mark as correct
- If multiple detections map to same ground truth, mark only one as correct
- **Precision** =  $\# \text{correct detections} / \text{total detections}$
- **Recall** =  $\# \text{ground truth with matched detections} / \text{total ground truth}$

# Tradeoff between precision and recall


- ML usually gives scores or probabilities, so we need to threshold
- Too low threshold → too many detections → low precision, high recall
- Too high threshold → too few detections → high precision, low recall
- Right tradeoff depends on application
  - Detecting cancer cells in tissue: need high recall
  - Detecting edible mushrooms in forest: need high precision

# Average precision





# Average precision



# *Average average precision*


- AP marks detections with overlap  $> 50\%$  as correct
- But may need better localization
- *Average* AP across multiple overlap thresholds
- Confusingly, still called average precision
- Introduced in COCO

# Mean and category-wise AP

- Every category evaluated independently
- Typically report mean AP averaged over all categories
- Confusingly called “mean Average Precision”, or “mAP”

# Why is detection hard(er)?

- Precise localization



# Why is detection hard(er)?

- Much larger impact of pose






# Why is detection hard(er)?

- Occlusion makes localization difficult



# Why is detection hard(er)?

- Counting



# Why is detection hard(er)?

- Small objects



# Object detection and localization

# Classification + Localization: Task

**Classification:** C classes

**Input:** Image

**Output:** Class label

**Evaluation metric:** Accuracy



CAT

**Localization:**

**Input:** Image

**Output:** Box in the image (x, y, w, h)

**Evaluation metric:** Intersection over Union



(x, y, w, h)

**Classification + Localization:** Do both

# Idea #1: Localization as Regression

**Input:** image



Only one object,  
simpler than detection

Neural Net



**Output:**

Box coordinates  
(4 numbers)

**Correct output:**  
box coordinates  
(4 numbers)

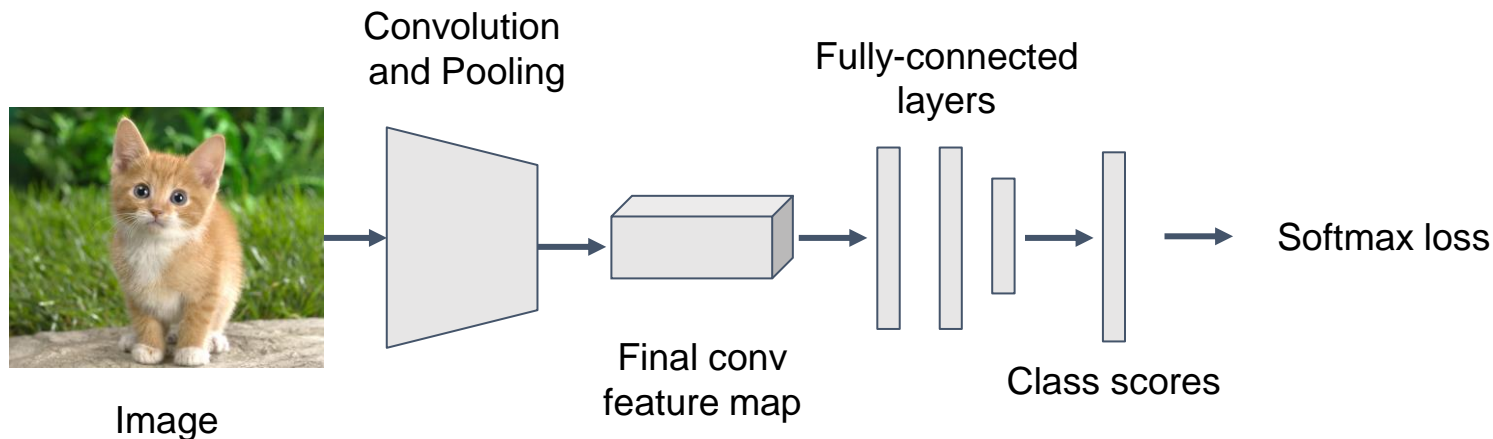


**Loss:**

L2 distance

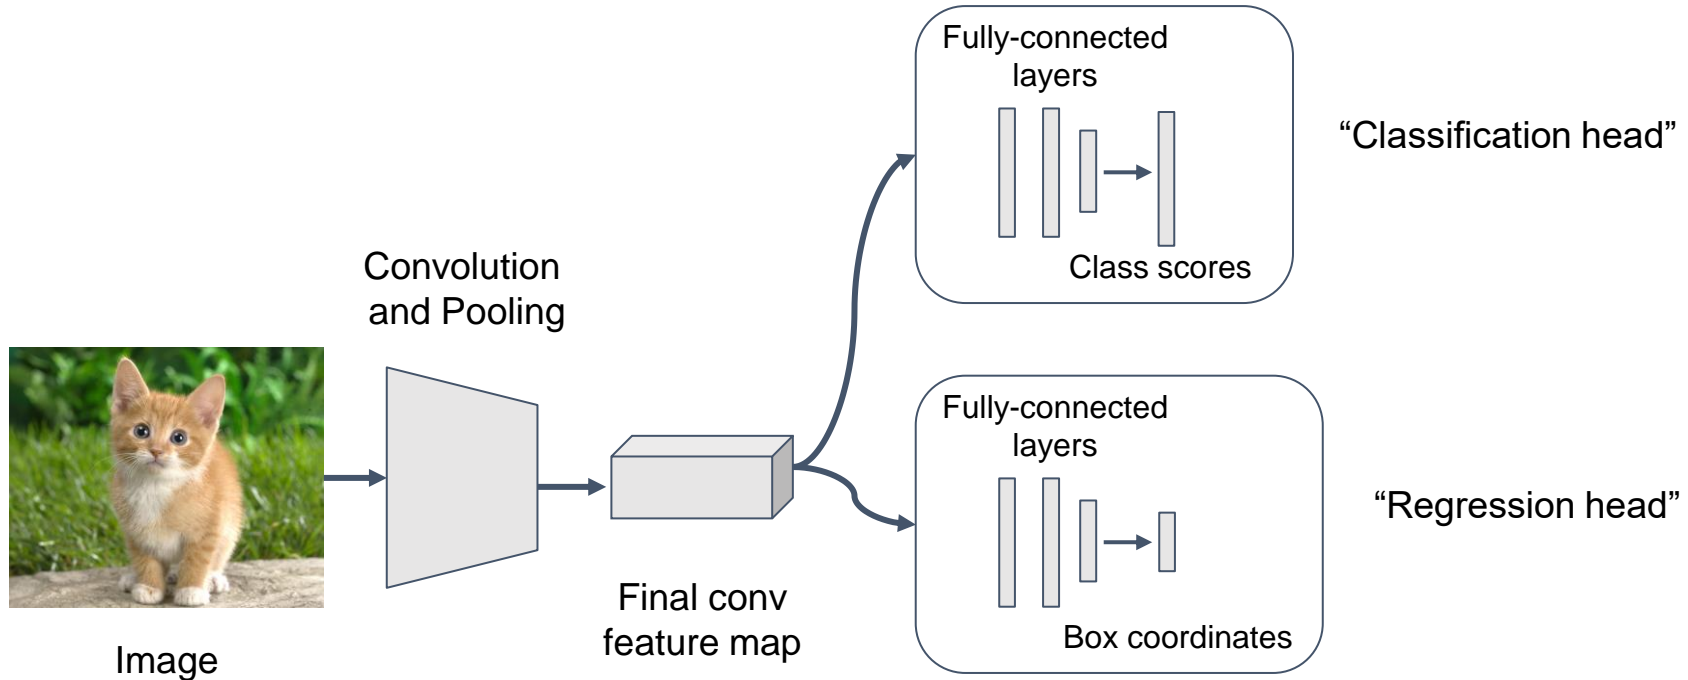
# Simple Recipe for Classification + Localization

**Step 1:** Train (or download) a classification model (AlexNet, VGG, GoogLeNet)



# Simple Recipe for Classification + Localization


**Step 2:** Attach a new fully-connected “regression head” to the network





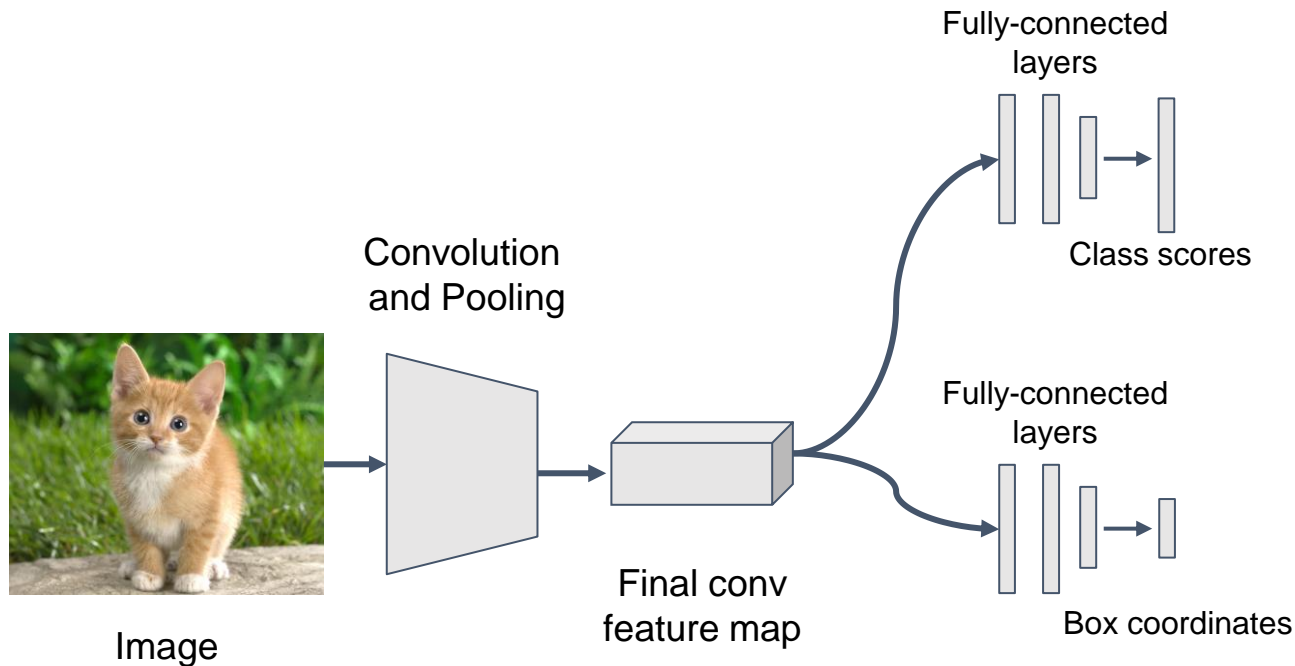
# Simple Recipe for Classification + Localization

**Step 3:** Train the regression head only with stochastic gradient descent (SGD) and L2 loss



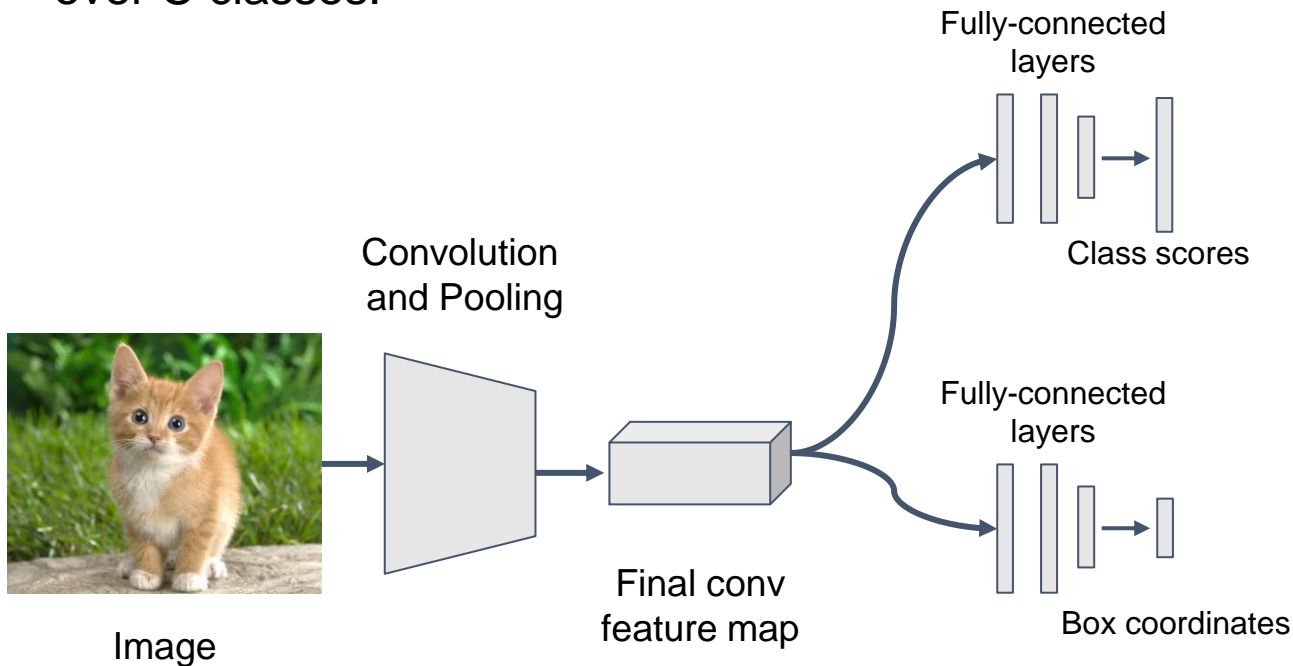
# Simple Recipe for Classification + Localization

**Step 4:** At test time use both heads

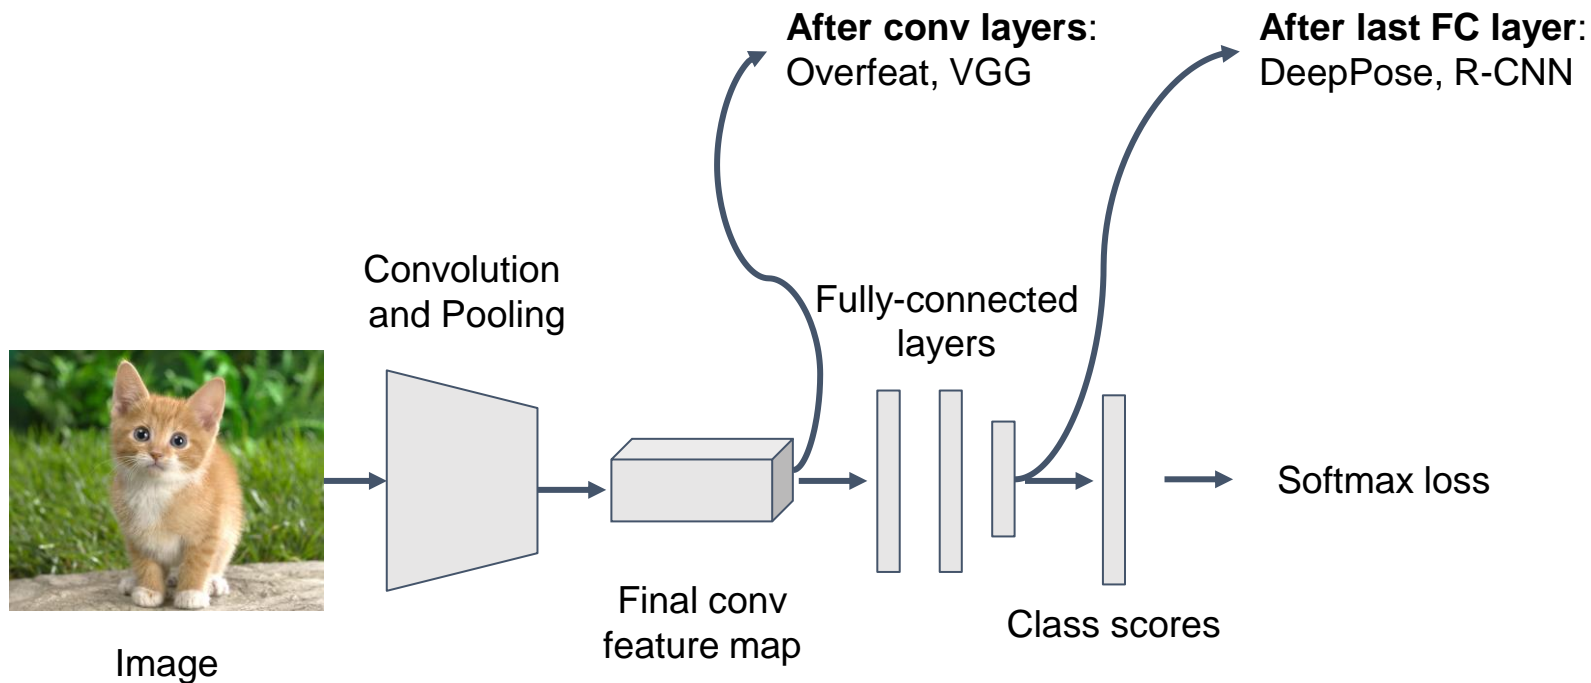


# Per-class vs class agnostic regression

Assume classification  
over  $C$  classes:

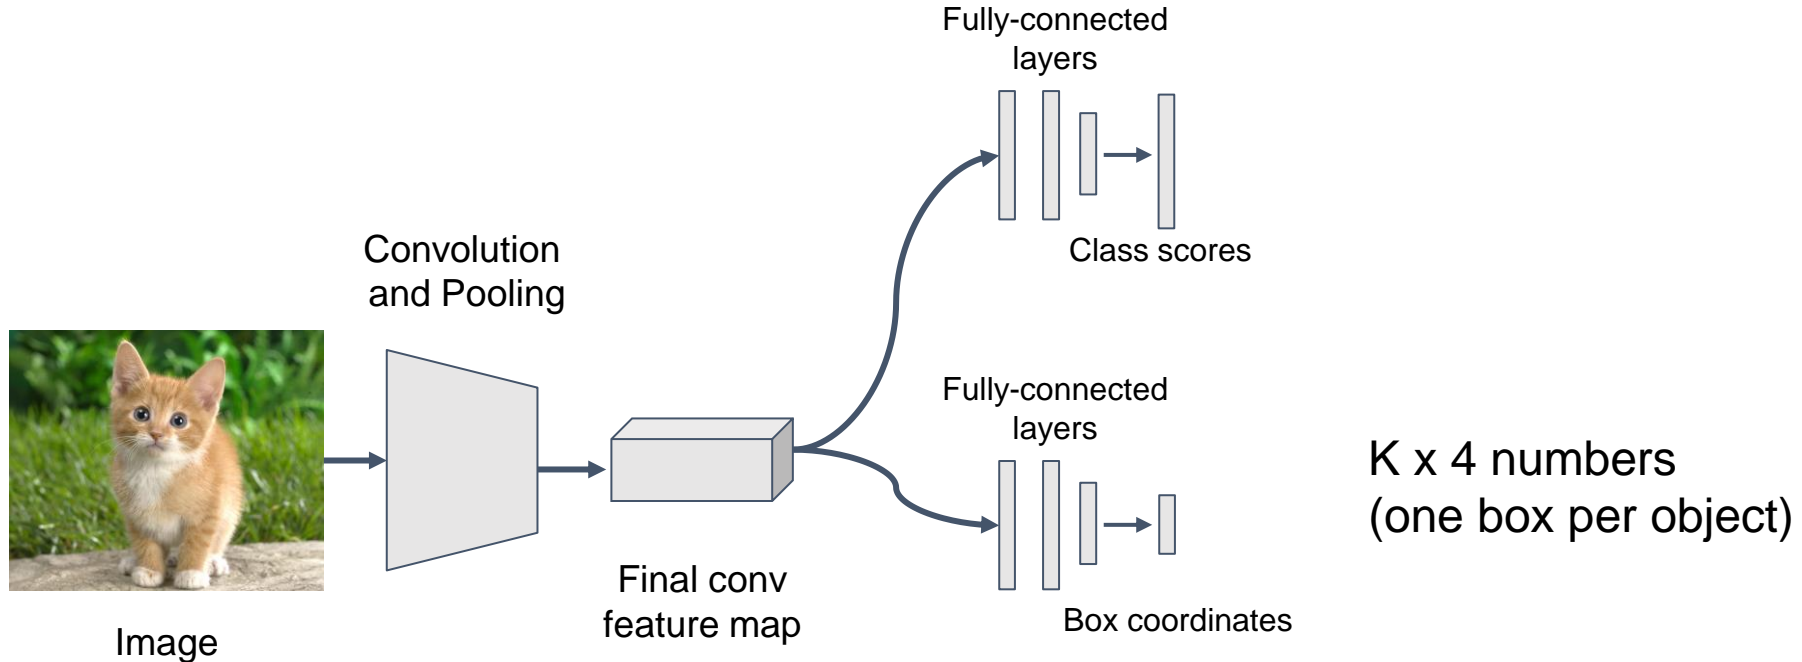


# Where to attach the regression head?




## Aside: Localizing multiple objects

Want to localize **exactly**  $K$  objects in each image (e.g. whole cat, cat head, cat's left ear, cat 's ear for  $K=4$ )



# Object Detection: Multiple Objects


Apply a classifier to many different crops of the image; the classifier classifies each crop as object or background



Dog? NO  
Cat? NO  
Background? YES

# Object Detection: Multiple Objects


Apply a CNN to many different crops of the image, CNN classifies each crop as object or background



Dog? YES  
Cat? NO  
Background? NO

# Object Detection: Multiple Objects

Apply a CNN to many different crops of the image, CNN classifies each crop as object or background




Dog? YES  
Cat? NO  
Background? NO



# Object Detection: Multiple Objects

Apply a CNN to many different crops of the image, CNN classifies each crop as object or background





Dog? NO  
Cat? YES  
Background? NO

Q: What's the problem with this approach?

# Object Detection: Multiple Objects

Apply a CNN to many different crops of the image, CNN classifies each crop as object or background



Dog? NO  
Cat? YES  
Background? NO


Problem: Need to apply CNN to huge number of locations, scales, and aspect ratios, very computationally expensive!

## Aside: Human Pose Estimation

Represent a person by  $K$  joints

Regress  $(x, y)$  for each joint from last fully-connected layer of AlexNet

(Details: Normalized coordinates, iterative refinement)



## Idea #2: Sliding Window

- Run classification + regression network at multiple locations on a high-resolution image
- Convert fully-connected layers into convolutional layers for efficient computation
- Combine classifier and regressor predictions across all scales for final prediction



# Sliding Window: Overfeat

Winner of ILSVRC 2013  
localization challenge



Image:  
3 x 221 x 221

Convolution  
+ pooling



Feature map:  
1024 x 5 x 5

FC

FC

4096

4096

FC

FC

4096

1024

FC

FC


Class scores:  
1000

Softmax  
loss

Euclidean  
loss

Boxes:  
1000 x 4

# Sliding Window: Overfeat




Network input:  
3 x 221 x 221




Larger image:  
3 x 257 x 257

# Sliding Window: Overfeat



Network input:  
 $3 \times 221 \times 221$




Larger image:  
 $3 \times 257 \times 257$


0.5	

Classification scores:  
 $P(\text{cat})$

# Sliding Window: Overfeat



Network input:  
 $3 \times 221 \times 221$




Larger image:  
 $3 \times 257 \times 257$

0.5	0.75


Classification scores:  
 $P(\text{cat})$



# Sliding Window: Overfeat



Network input:  
 $3 \times 221 \times 221$




Larger image:  
 $3 \times 257 \times 257$


0.5	0.75
0.6	

Classification scores:  
 $P(\text{cat})$

# Sliding Window: Overfeat



Network input:  
 $3 \times 221 \times 221$




Larger image:  
 $3 \times 257 \times 257$


0.5	0.75
0.6	0.8

Classification scores:  
 $P(\text{cat})$

# Sliding Window: Overfeat



Network input:  
 $3 \times 221 \times 221$




Larger image:  
 $3 \times 257 \times 257$

0.5	0.75
0.6	0.8


Classification scores:  
 $P(\text{cat})$

# Sliding Window: Overfeat

Greedily merge boxes and scores (details in paper)



Network input:  
3 x 221 x 221



Larger image:  
3 x 257 x 257


0.8

Classification score:  
P(cat)


# Sliding Window: Overfeat

In practice use many sliding window locations and multiple scales

Window positions + score maps



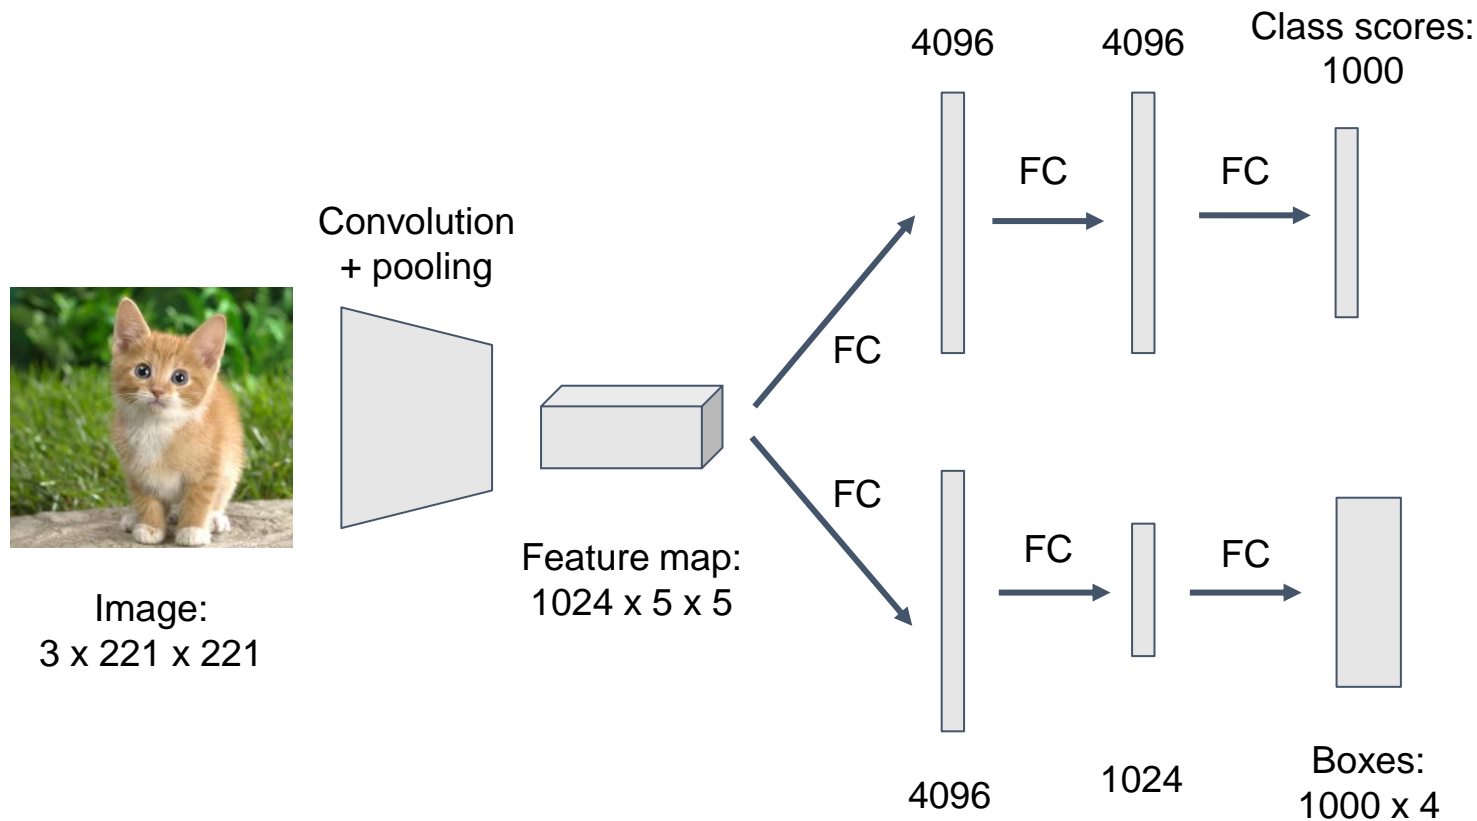
Box regression outputs



Final Predictions



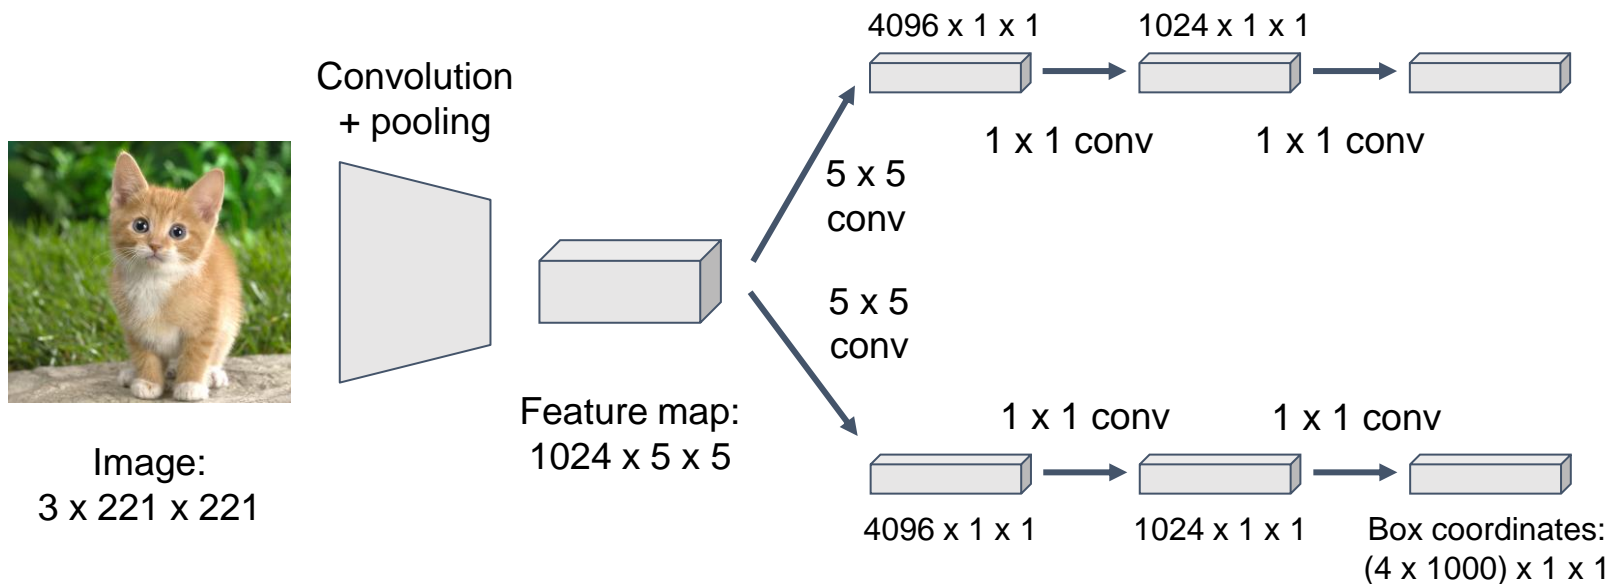
# Efficient Sliding Window: Overfeat



# Efficient Sliding Window: Overfeat


Efficient sliding window by converting fully-connected layers into convolutions

Class scores:  
 $1000 \times 1 \times 1$




# Efficient Sliding Window: Overfeat

**Training time:** Small image, 1 x 1 classifier output




**Test time:** Larger image, 2 x 2 classifier output, only extra compute at yellow regions





# ImageNet Classification + Localization



**AlexNet:** Localization method not published

**Overfeat:** Multiscale convolutional regression with box merging

**VGG:** Same as Overfeat, but fewer scales and locations; simpler method, gains all due to deeper features


**ResNet:** Different localization method (Region Proposal Network - RPN) and much deeper features

# “Sliding Window” Detection


background background background background background background background background



background person




# “Sliding Window” Detection



Compute within-region features,  
then classify


# “Sliding Window” Detection



Typical to enlarge region to include some “context”


# Sliding window placement

Slide over *fine grid*  
in x, y, scale, aspect ratio



**Slow and Accurate**

Slide over *coarse grid*  
in x, y, scale, aspect ratio



**Fast and Not-so-accurate**  
(... or can it be?)




# Bounding Box Regression



**Idea:**

**Also predict  
continuous offset  
from anchor to "snap"  
onto object**

# Bounding Box Regression

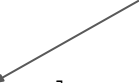


**Idea:**

**Also predict  
continuous offset  
from anchor to “snap”  
onto object**

# Typical Training Objective

## Per-anchor Loss:

$$L(\text{anchor } \mathbf{a}) = \alpha * \delta(\mathbf{a} \text{ has matching groundtruth}) * L_2(\mathbf{t}^{\text{loc}}, W^{\text{loc}} \cdot \mathbf{v}_{ij}) \\ + \beta * \text{SoftMaxCrossEntropy}(\mathbf{t}^{\text{cls}}, W^{\text{cls}} \cdot \mathbf{v}_{ij})$$


Common to use other location losses here...


**Total Loss:** Average per-anchor loss over anchors

**Challenge:** Dealing with class imbalance (usually way more negative anchors (class 0) than positive anchors)

**Solutions:** Subsampling negative anchors, downweighting the loss contribution of negatives, hard mining, etc...



# Dealing with multiple detections of the same object



Duplicate detection problem: Typically many anchors will detect the same underlying object and give slightly different boxes, with slightly different scores.

Solution: remove detections if they overlap too much with another higher scoring detection.


# Non Max Suppression (NMS)

**Algorithm:**

- 1. Sort detections in decreasing order with respect to score**
- 2. Iterate through sorted detections:**  
Reject a detection if it overlaps with a previous (unrejected) detection with IOU greater than some threshold
- 3. Return all unrejected detections**

**Some shortcomings of NMS to remember:**

- Imposes a hard limitation on how close objects can be in order to be detected**
- Similar classes do not suppress each other**

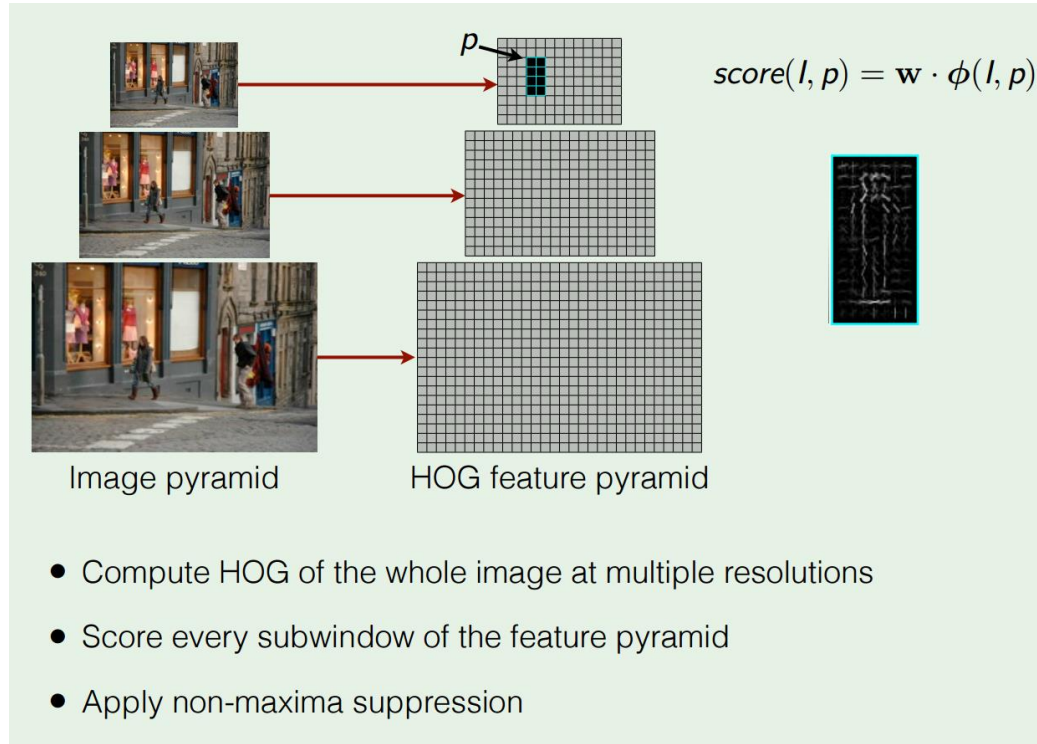


# Detection as Classification

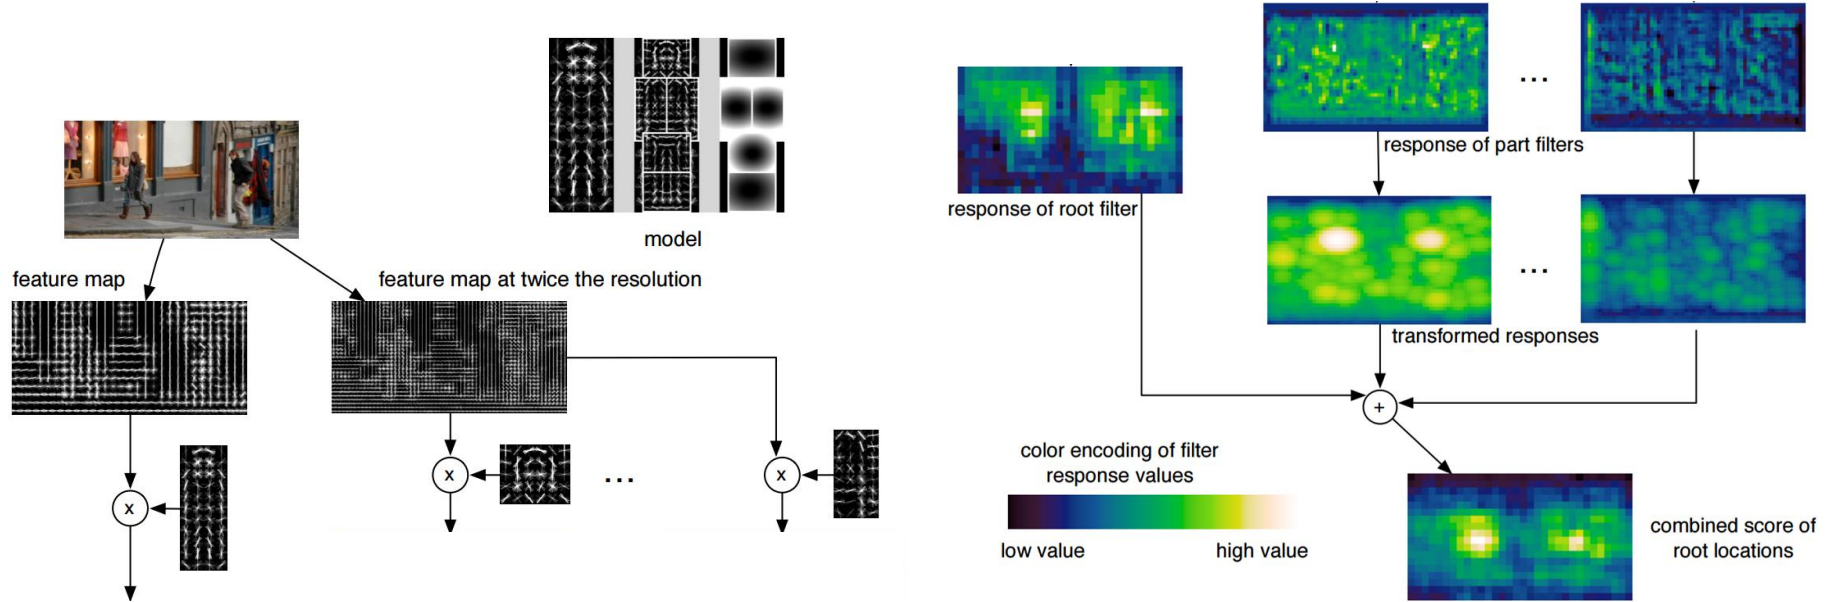
**Problem:** Need to test many positions and scales

**Solution:** If your classifier is fast enough, just do it

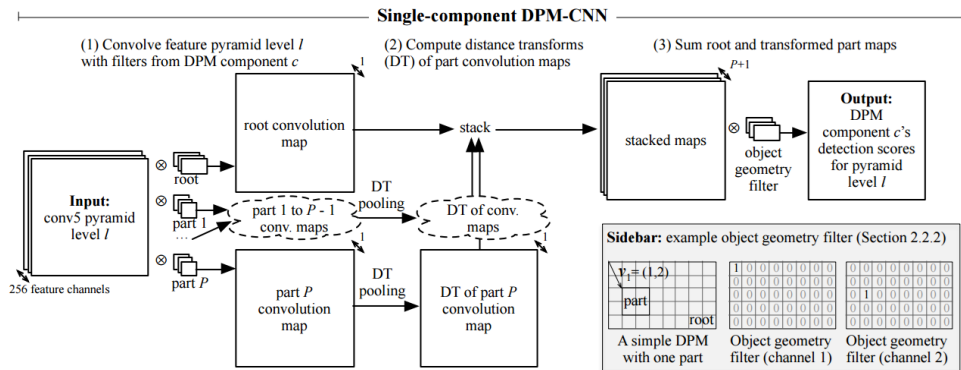
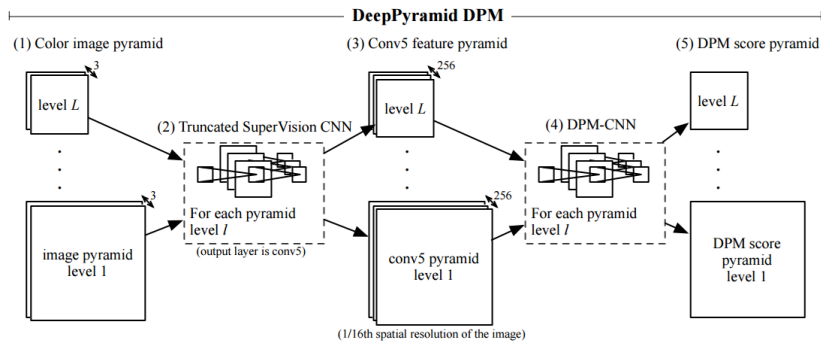
# Histogram of Oriented Gradients



# Deformable Parts Model (DPM)



## Aside: Deformable Parts Models are CNNs?



# Detection as Classification

**Problem:** Need to test many positions and scales,  
and use a computationally demanding classifier (CNN)

**Solution:** Only look at a tiny subset of possible positions