



Part VII: Classification

Nicola Conci
nicola.conci@unitn.it

Pattern recognition problems



- When you have to **make a decision** about a content of an image
- When you have to determine **what the object** present in the picture **represents**
- Need to use the **features of the appearance** of the object to determine the category it belongs to
- Examples:
 - In a supermarket implement a system that recognizes vegetables
 - At the gate of your house, restrict the access to specific subjects or cars
 - At the entrance of a parking lot determine the category of the entering vehicle
 - On a robot, look for objects (doors, humans, obstacles)

Classes




- Objects, patterns, textures are arranged in classes
- A **class is made up of descriptions** provided by a number of examples
- Classes contain items that share **similar properties**
- Goal of a classifier is to take an object as input and to output the class label it belongs to

The classification process



- Starting from an input image the steps are:
 - Segment the image (if necessary)
 - Extract the features
 - Use the feature vector to feed a classifier
 - Output a label



Basic principles



- Children learn new things by examples
- In order to distinguish a dog from a cat, they rely on observing samples of dogs and cats
- We realize they have learned the concept by the time they see a new sample and they're able to recognize it
- Teaching a computer works the same way

Subject to failures




- Classification is a tough job
- It usually requires **A LOT** of annotated data
- Machines may take days or weeks in order to learn concepts
- Still they are likely to fail due to:
 - Lack of data
 - Wrong annotations
 - Occlusions
 - Perspective
 - Similarity among classes

Examples of Failures...



... and also very good results



The pipeline



Define the task


e.g. Dogs vs cats




Collect data



Design features



Train model




Test model

Simple problem: regression



- Use case: predict the share price of a company that is about to go public
- Training data coming from similar companies
- Features: companies revenue and share price



- Define loss function (our score to be minimized)
 - Find the parameters of the line (slope, intercept)
 - Use the line for testing
- In our problem:
- We use the training to determine the line
 - In testing, given the revenue we can predict the share price

More formally



- We have a set of training samples $D = \{Z_1, Z_2, \dots, Z_n\}$
- We have an unknown process $P(Z)$
- And a loss function $L(f, Z)$ where f is the decision function
- In **supervised** learning each example is a pair $Z(X, Y)$ where
 - X contains the features for that sample
 - Y is the known/expected output
 - f takes X as an argument and outputs something in the range of Y
- Loss function defined as:

$$L(f, (X, Y)) = \|f(X) - Y\|^2$$

In case of multi-class



- Y is a class (finite integer) and the loss function is the negative conditional log-likelihood
- $f_i(X)$ estimates $P(Y=i | X)$

$$L(f, (X, Y)) = -\log f_Y(X)$$

- With the constraint that

$$f_Y(X) \geq 0, \sum_i f_i(X) = 1$$

Unsupervised learning



- Learn f to characterize $P(Z)$
- Through clustering the space is partitioned in regions centered around a prototype (centroid)
 - K-Means (hard partitioning)
 - GMM (soft partitioning), each Z has a probability



Classification output - evaluation

- As an output you may have:
 - Correct detection
 - The item is associated to the correct class
 - Wrong detection
 - An item is associated to the wrong class
- In a binary classification problem:
 - True positive (hit)
 - True negative
 - False positive
 - False negative (miss)

Reject option



- As we have seen the classification is subject to errors
- Errors tend to arise in the presence of those samples for which largest of the posterior probability $P(Y=i|X) \ll 1$
→ *intuitively, this means that the class with the highest likelihood is still very low.*
- In such cases the system can implement a so-called *reject option*, or *reject class*
- This helps avoiding taking decisions for those samples that are difficult to classify

A simple example




- Let us take as a very simple example a binary classification problem
- The output can only be true or false
- Ex: Does a patient have a disease?

| | | Current value | |
|-------------|---|---------------|----|
| | | p | n |
| Test Result | p | TP | FP |
| | n | FN | TN |

The ROC curve



- Receiver Operating Characteristic
- Plot of TPR vs FPR
- Each setup of the system is a point in the ROC space



Precision and Recall



- The system is usually targeted at finding elements that satisfy a query
 - **Precision:** $TP / (TP + FP) = \text{hit} / (\text{all retrieved})$
 - **Recall:** $TP / (TP + FN) = \text{hit} / (\text{hit} + \text{miss})$
- In terms of probability it means:
 - Precision: probability that a randomly selected elements is relevant
 - Recall: probability that a randomly relevant element is retrieved in the search

Precision and Recall in practice



- Assume we want to find all pictures that contain faces
- We have 200 images and only 100 contain faces
- The classifier returns only 75 (TP) out of 100 but also additional 30 images (FP).
- Precision is $75/(75+30) = 71\%$
- Recall is $75/100 = 75\%$

The confusion matrix



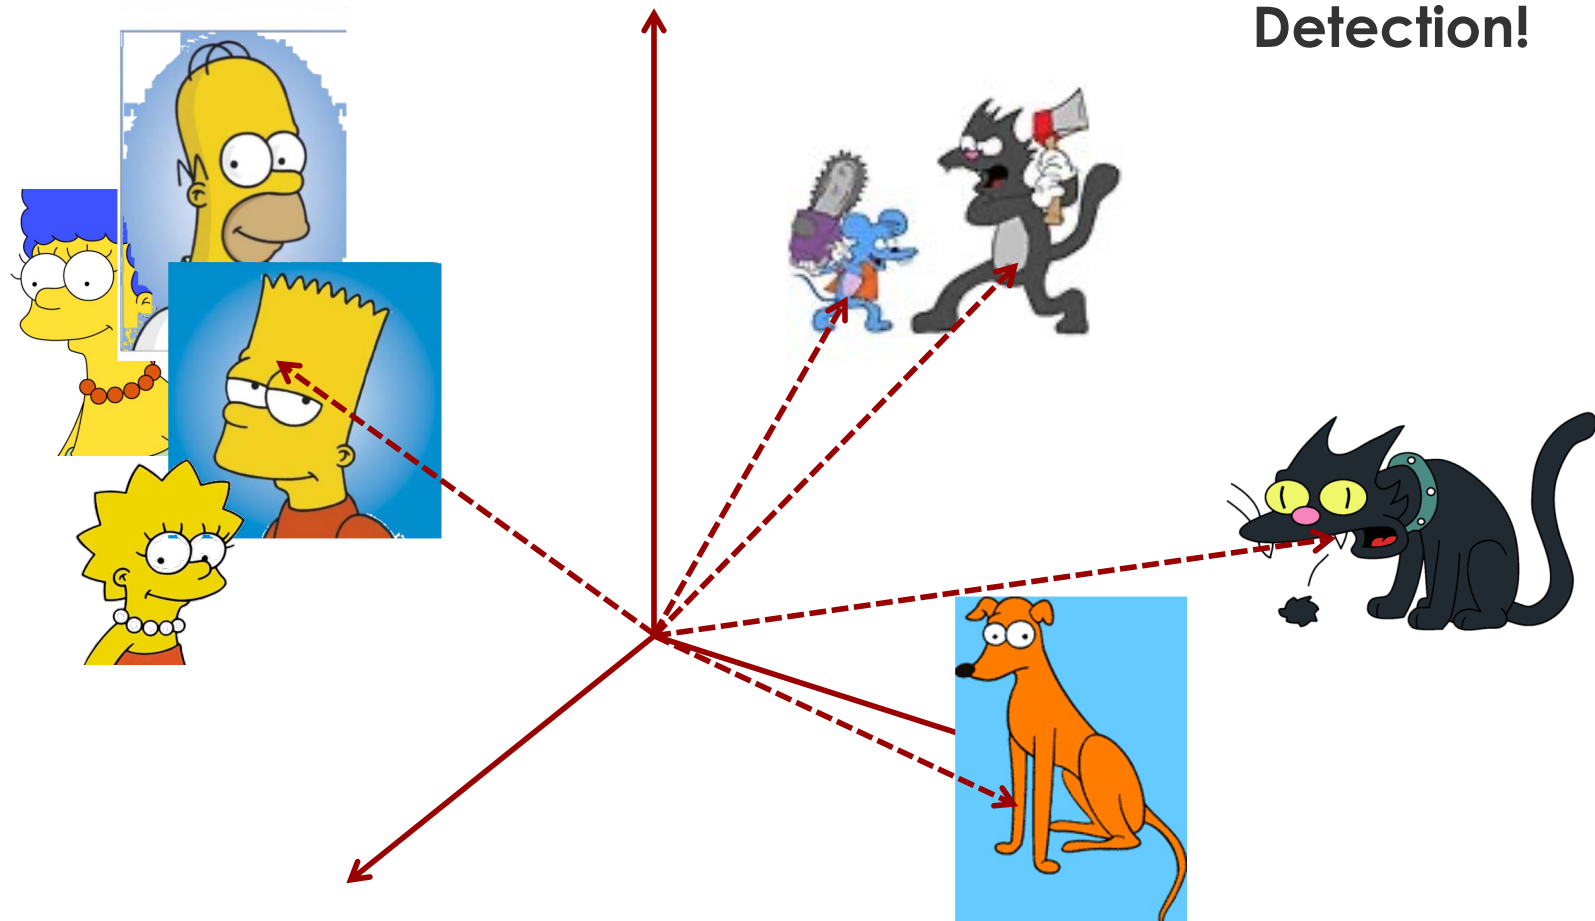
- Used in case we have multiple classes
- In vegetables: Potatoes – Lettuce –Tomatoes
- The simple true/false matrix is not sufficient anymore
- Construct a matrix with the correct classification on the diagonal and the false positives in the other cells

The face detection problem




- Once I have identified the significant features for the object that I'm looking for, I can take some samples and train the classifier
- If I'm looking for faces the problem can be:
 - Find all faces (binary - detection)
 - Find Bart – Lisa – Homer – Marge (recognition)
- The two scenarios are very different and involve the application of different algorithms

The classification problem



The classification problem



The Viola-Jones algorithm




- The Viola-Jones algorithm probably the most widespread face detector
- Available in OpenCV
- Goal:
 - Implement a robust classifier using simple features
 - Based on binary features

Features




- 14 Haar-like features
- Pretty simple features: vertical, horizontal, and diagonal


1. Edge features



2. Line features



3. Center-surround features



Features



- Two/Three/Four-rectangle features
- The sum of the pixels within the white rectangles are subtracted from the sum of pixels in the black rectangles
- Rectangle features can be computed easily using an intermediate representation for the image called the *integral image*


$$ii(x, y) = \sum_{x' \leq x, y' \leq y} i(x', y')$$

- The *integral image* contains the sum of the above and left pixels w.r.t the current location $i(x, y)$

The integral image



- The sum of the values in D can be computed using four reference points:
- The value in 2 is $A+B$
- At 3 it is $A+C$
- At 4 $A+B+C+D$
- Sum in D is $4+1-(2+3)$



Recursions



- Using the following recursions, the integral image can be computed over the entire image in one single pass

$$s(x, y) = s(x, y - 1) + i(x, y)$$

$$ii(x, y) = ii(x - 1, y) + s(x, y)$$

- s is the cumulative row sum

Classifiers



- In literature a lot of different classification algorithms have been proposed
 - Artificial neural networks
 - Feedforward
 - Radial Basis Function
 - Multi Layer Perceptron
 - Support Vector Machines

- **Boosting**

AdaBoost



- It consists of the implementation of a strong classifier that relies on a combination of weak classifiers

$$f(x) = \sum_{t=1}^T \alpha_t h_t(x)$$

- $h(x)$ is the weak classifier, a feature evaluated in the image
- $H(x) = \text{sign}(f(x))$ is the strong classifier



$$h_i(x) = \begin{cases} 1 & \text{if } f_i > \text{threshold} \\ -1 & \text{if } f_i < \text{threshold} \end{cases}$$

$$f_i = \text{Sum}(r_{i, \text{white}}) - \text{Sum}(r_{i, \text{black}})$$

AdaBoost



- Given a set of points (\mathbf{x}_i, y_i) , $i=1 \dots m$, with $y=\pm 1$
- Initialize weights $D_1(i) = 1/m$
- Evaluate the feature
- Find
$$h_t = \operatorname{argmin}_{h_j \in H} \varepsilon_j = \sum_{j=1}^m D_t(i) I(y_i \neq h_j(x_i))$$
- I is the *Indicator function* (binary 1 or 0)
- If error > 0.5 stop
- Choose α_t

It represents the sum of the weights of the misclassified samples

$$\alpha_t = \frac{1}{2} \ln \left(\frac{1 - \varepsilon_t}{\varepsilon_t} \right)$$

AdaBoost



- Update weights

$$D_{(t+1)}(i) = \frac{D_t(i)e^{-\alpha_t y_i h_t(x_i)}}{Z_t}$$

Z is a
normalization
factor

$$Z_t = \sum_i D_t(i) e^{-\alpha_t y_i h_t(x_i)}$$

Samples that are
more difficult to
classify will have
higher weight in the
next iteration

$$-\alpha_t y_i h_t(x_i) \begin{cases} < 0, y(i) = h_t(x_i) \\ > 0, y(i) \neq h_t(x_i) \end{cases}$$


- The final output of the classifier becomes:

$$H(x) = \text{sign} \left(\sum_{t=1}^T \alpha_t h_t(x) \right)$$

AdaBoost




- For example



Cascade of classifiers




- In AdaBoost the combination of the weak classifiers improves the speed of the classification process
- In the algorithm proposed by Viola and Jones, classifiers are used in cascade, further reducing the computational complexity



Cascade of classifiers



- Goal: quickly remove false negatives and focus on the positive samples
- In face detection, basic features can be used to exclude non-faces



- The output of the first classifier is used to trigger the second one
- Each stage is trained using AdaBoost
- At each stage the negatives are rejected

Training and testing



- In the cascade, binary features are evaluated on training images of equal size (24x24pix)
- Classification:
 - During the training stage, the boundaries of the classifiers are learned
 - During the testing phase, the learned classifiers are used to evaluate *unknown* samples
- Images are not scaled at 24x24, so a multi-scale analysis must be conducted
- Scaling is performed at the feature level, not image → fast by considering the type of detector [binary + integral image]
- In case of multiple detections at different scales, windows are averaged

Complexity and concluding remarks



- Complexity depends on :
 - the number of classifiers in the cascade (mainly in training)
 - The number of levels in the multi-scale analysis
 - The sampling distance between the windows
- Images used for training/testing could be acquired in different environmental conditions, so light is in general different
- A normalization would be highly desirable
- Here variance normalization in the sub-window

$$\sigma^2 = m^2 - \frac{1}{M} \sum_{x \in M} x^2$$